



Jaeger Project Deep Dive

Annanay Agarwal (Grafana)
Pavol Loffay (Traceable.ai)
Yuri Shkuro (Facebook)

KubeCon + CloudNativeCon NA 2020 Virtual
Thu, Nov 19 • 2:55 pm - 3:30 pm

About

- Yuri Shkuro (<https://github.com/yurishkuro>)
 - Software engineer
 - Maintainer of Jaeger, OpenTracing, OpenTelemetry
 - Author of “[Mastering Distributed Tracing](#)” book

Agenda

- Observability and tracing
- Jaeger features
- Jaeger architecture
- Sampling
- Jaeger and OpenTelemetry
- Jaeger on Kubernetes

About

- Annanay Agarwal (<https://github.com/annanay25>)
 - Software developer at Grafana Labs
 - Contributor to Jaeger and OpenTelemetry projects

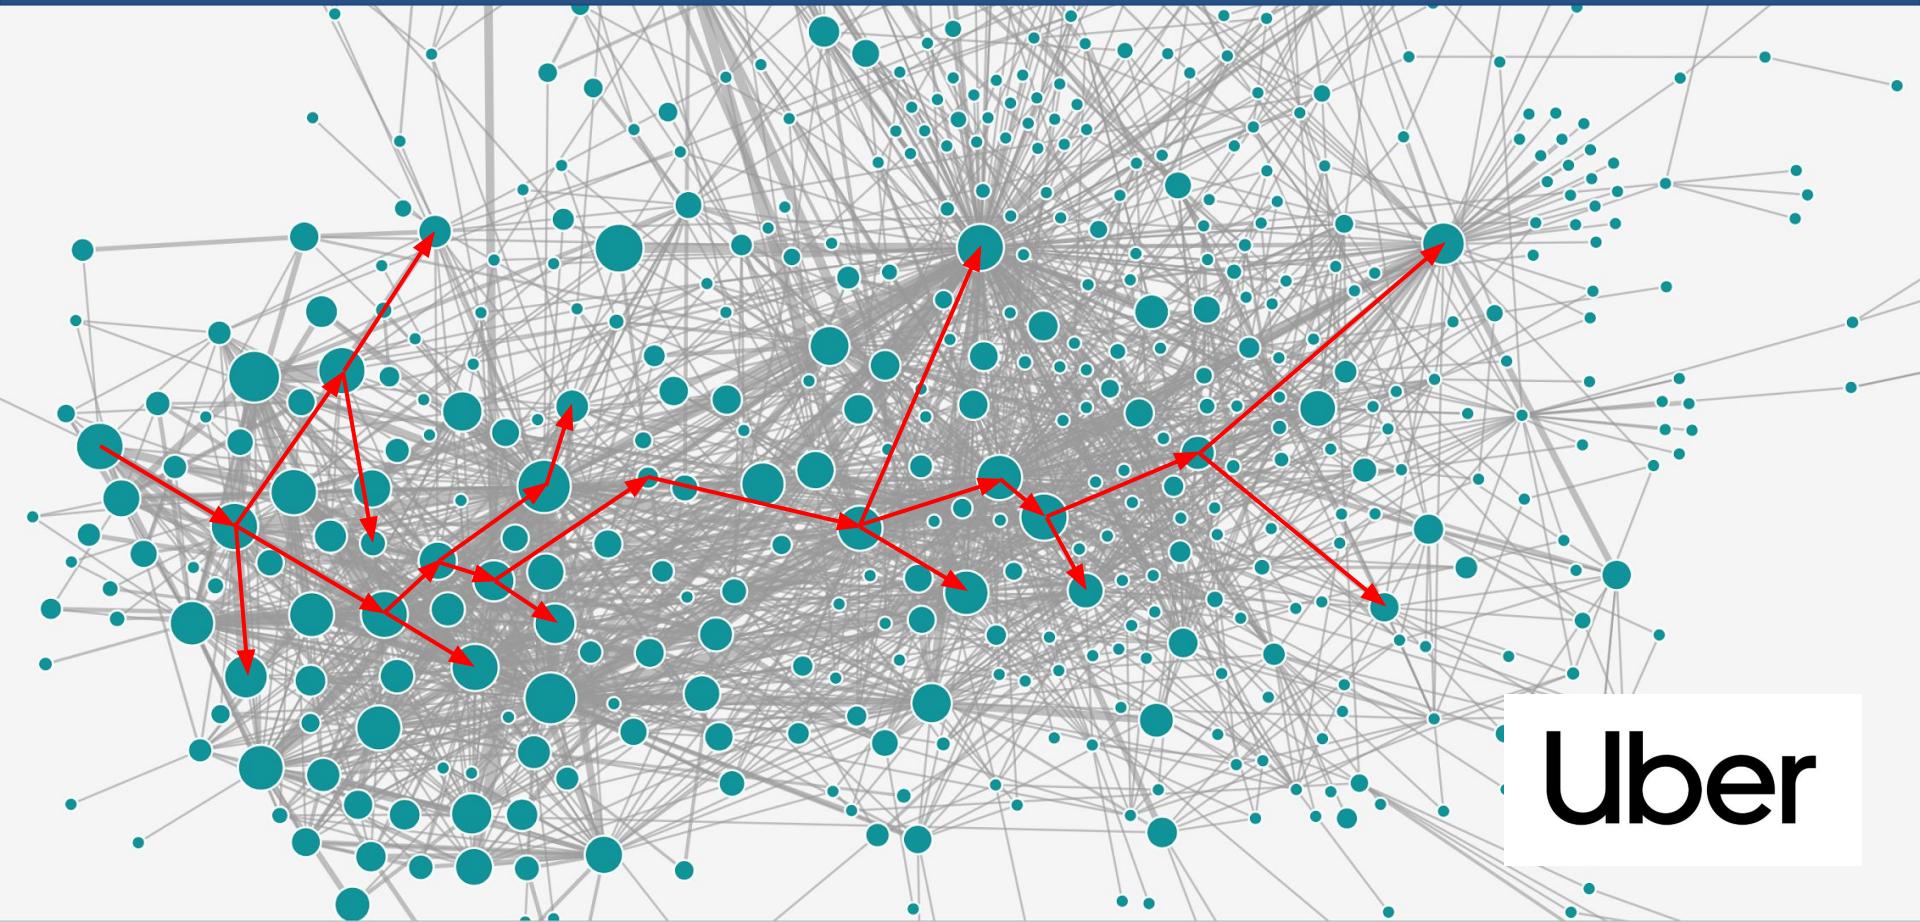


CLOUD NATIVE
COMPUTING FOUNDATION

What is Tracing & Why?

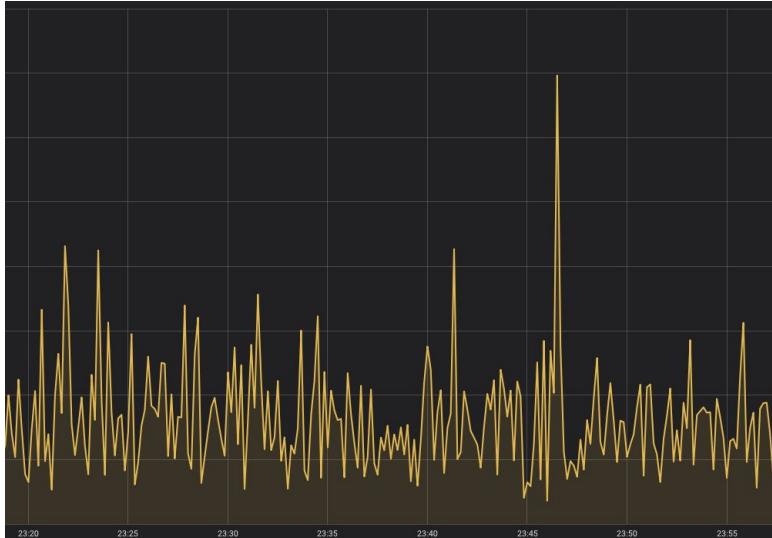
Concepts and terminology

BILLIONS of times a day!



Uber

Metrics



```
http_request_duration_sec{"app=ice-cream-shop"} 10s
```

Metrics - Cardinality

```
http_request_duration_sec{"app=ice-cream-shop"  
datacenter="us-central", env="production",  
service="cart-manager", path="/api/order",  
func_name="my-func"} 6s
```



Logs - stack trace?

```
2017-05-05T17:56:10.467+0000  INFO log/spanlogger.go:46  HTTP {"service": "frontend", "method": "GET", "url": "/favicon.ico"}  
2017-05-05T17:59:38.356+0000  INFO log/spanlogger.go:46  HTTP {"service": "frontend", "method": "GET", "url": "/favicon.ico"}  
2017-05-05T17:59:38.463+0000  INFO log/spanlogger.go:46  HTTP {"service": "frontend", "method": "GET", "url": "/favicon.ico"}  
2017-05-05T18:00:03.975+0000  ERROR log/spanlogger.go:46  redis timeout {"service": "customer", "error": "redis timeout", "method": "GET", "url": "/customer?customer_id=123&nonce=0.8534872559455979"}  
2017-05-04T08:36:14.417+0000  INFO log/spanlogger.go:46  Loading customer {"service": "customer", "component": "customer_client", "customer_id": "123"}  
2017-05-04T08:36:14.418+0000  INFO log/spanlogger.go:46  Loading customer {"service": "customer", "component": "customer", "customer_id": "123"}  
2017-05-04T08:36:14.419+0000  INFO log/spanlogger.go:46  Loading customer {"service": "customer", "component": "mysql", "customer_id": "123"}  
2017-05-04T08:36:14.419+0000  INFO log/spanlogger.go:46  Finding nearest drivers {"service": "driver", "location": "115.2777"}  
2017-05-04T08:36:14.729+0000  INFO log/spanlogger.go:46  Finding for nearby drivers {"service": "driver", "location": "115.2777"}  
2017-05-04T08:36:14.729+0000  INFO log/spanlogger.go:46  {"method": "GET", "url": "/nearest_driver?customer_id=123&redis_timeout"}  
github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap.Stats  
github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap.(*Field).Field  
github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap.field.go:209  
github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap.(*Logger).Check  
{/users/yurishkun/golang/src/github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap/logger.go#L273}  
github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap.(*Field).Field  
{/users/yurishkun/golang/src/github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap/logger.go#L176}  
github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap.(*Logger).Error  
{/users/yurishkun/golang/src/github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap/logger.go#45}  
github.com/uber/jaeger/examples/hotrod/pkg/log/spanlogger.go:45  
github.com/uber/jaeger/examples/hotrod/vendor/github.com/uber/channeldriver.(*Driver).GetDriver  
autogenerated#19  
github.com/uber/jaeger/examples/hotrod/vendor/github.com/uber/channeldriver.(*Driver).GetDriver  
github.com/uber/jaeger/examples/hotrod/services/drivers.(*Server).FindNearest  
github.com/uber/jaeger/examples/hotrod/services/drivers.(*Server).FindNearest  
github.com/uber/jaeger/examples/hotrod/services/drivers.(*Driver).FindNearest  
{/users/yurishkun/golang/src/github.com/uber/jaeger/examples/hotrod/services/drivers/nearest.go#19}  
github.com/uber/jaeger/examples/hotrod/service/drive/thrift/gen/processor.(*ChannelServer).HandleInNearest  
{/users/yurishkun/golang/src/github.com/uber/jaeger/examples/hotrod/service/drive/thrift/gen/processor/channel.go#92}  
github.com/uber/jaeger/examples/hotrod/service/drive/thrift/gen/processor.(*ChannelServer).HandleInNearest  
{/users/yurishkun/golang/src/github.com/uber/jaeger/examples/hotrod/service/drive/thrift/gen/processor/channel.go#76}  
github.com/uber/jaeger/examples/hotrod/service/drive/thrift/gen/processor.(*ChannelServer).HandleInNearest  
{/users/yurishkun/golang/src/github.com/uber/jaeger/examples/hotrod/service/drive/thrift/gen/processor/channel.go#133}  
github.com/uber/jaeger/examples/hotrod/service/drive/thrift/gen/processor.(*ChannelServer).HandleInNearest  
{/users/yurishkun/golang/src/github.com/uber/jaeger/examples/hotrod/service/drive/thrift/gen/processor/channel.go#293}  
github.com/uber/jaeger/examples/hotrod/service/drive/thrift/gen/processor.(*ChannelServer).HandleInNearest  
{/users/yurishkun/golang/src/github.com/uber/jaeger/examples/hotrod/service/drive/thrift/gen/processor/channel.go#118}  
github.com/uber/jaeger/examples/hotrod/vendor/github.com/uber/channeldriver.(*Channel).GetHandlers  
github.com/uber/jaeger/examples/hotrod/vendor/github.com/uber/channeldriver.(*Channel).GetHandlers  
github.com/uber/jaeger/examples/hotrod/vendor/github.com/uber/channeldriver.(*Channel).DispatchInbound  
github.com/uber/jaeger/examples/hotrod/vendor/github.com/uber/channeldriver.(*Channel).DispatchInbound  
autogenerated#19  
github.com/uber/jaeger/examples/hotrod/vendor/github.com/uber/channeldriver.(*Driver).FindNearest  
github.com/uber/jaeger/examples/hotrod/vendor/github.com/uber/channeldriver.(*Driver).FindNearest  
github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap.Stats  
github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap.field.go:209  
github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap.(*Logger).Check  
{/users/yurishkun/golang/src/github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap/logger.go#L273}  
github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap.(*Logger).Error  
{/users/yurishkun/golang/src/github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap/logger.go#176}  
github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap.(*Field).Field  
{/users/yurishkun/golang/src/github.com/uber/jaeger/examples/hotrod/vendor/go.uber.org/rap.field.go#45}  
github.com/uber/jaeger/examples/hotrod/vendor/github.com/uber/channeldriver.(*Driver).FindNearest  
{/users/yurishkun/golang/src/github.com/uber/jaeger/examples/hotrod/services/drivers/nearest.go#19}
```

Logs are a mess: concurrent requests, multiple hosts, impossible to correlate.

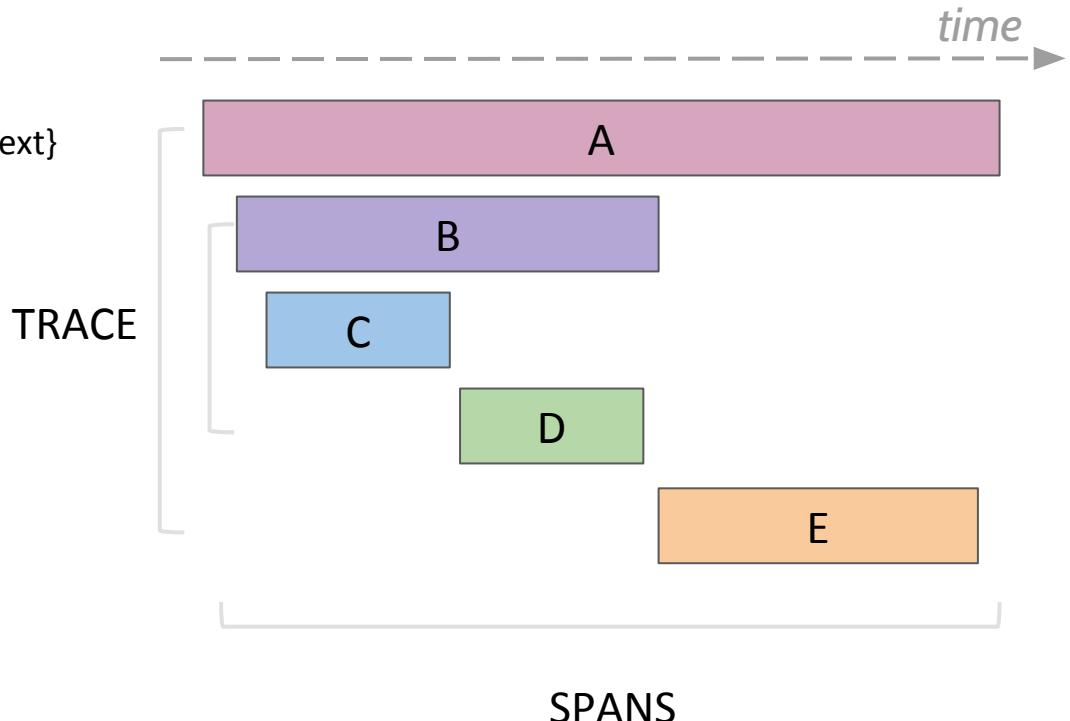
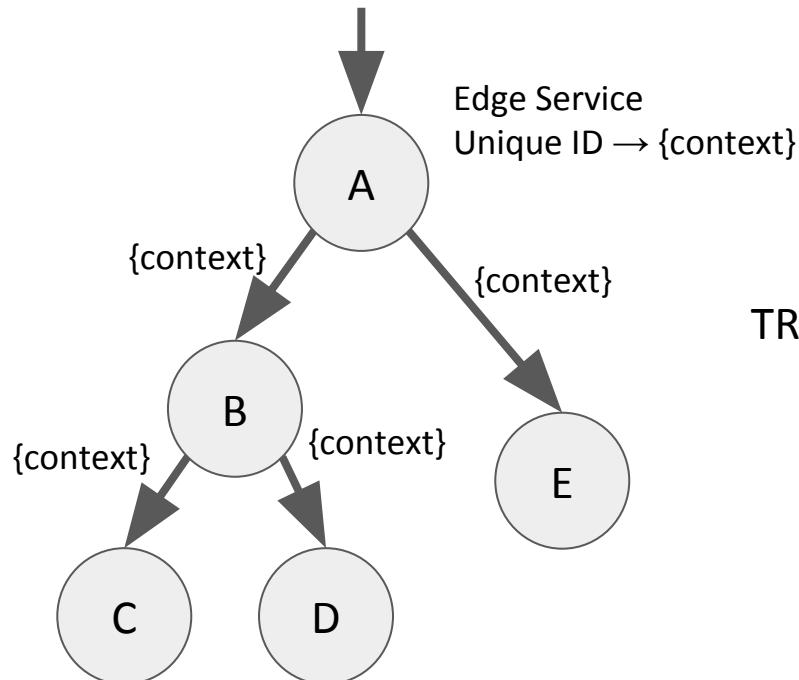
Monitoring tools must tell stories!

Do you like debugging
without a stack trace?

We need to monitor
distributed transactions
⇒ **distributed tracing!**



Context Propagation & Distributed Tracing





Let's look at some traces

<http://bit.do/jaeger-hotrod>



Service dependencies diagram

Jaeger UI

Lookup by Trace ID...

Search

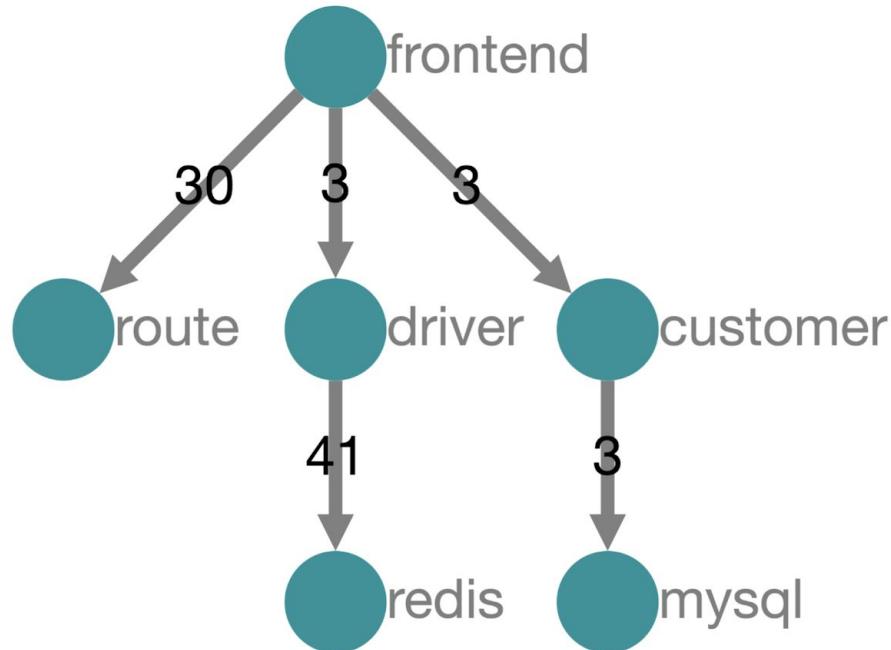
Compare

Dependencies

About Jaeger

Force Directed Graph

DAG



Transitive Service Graphs

4 Traces

Sort: Most Recent

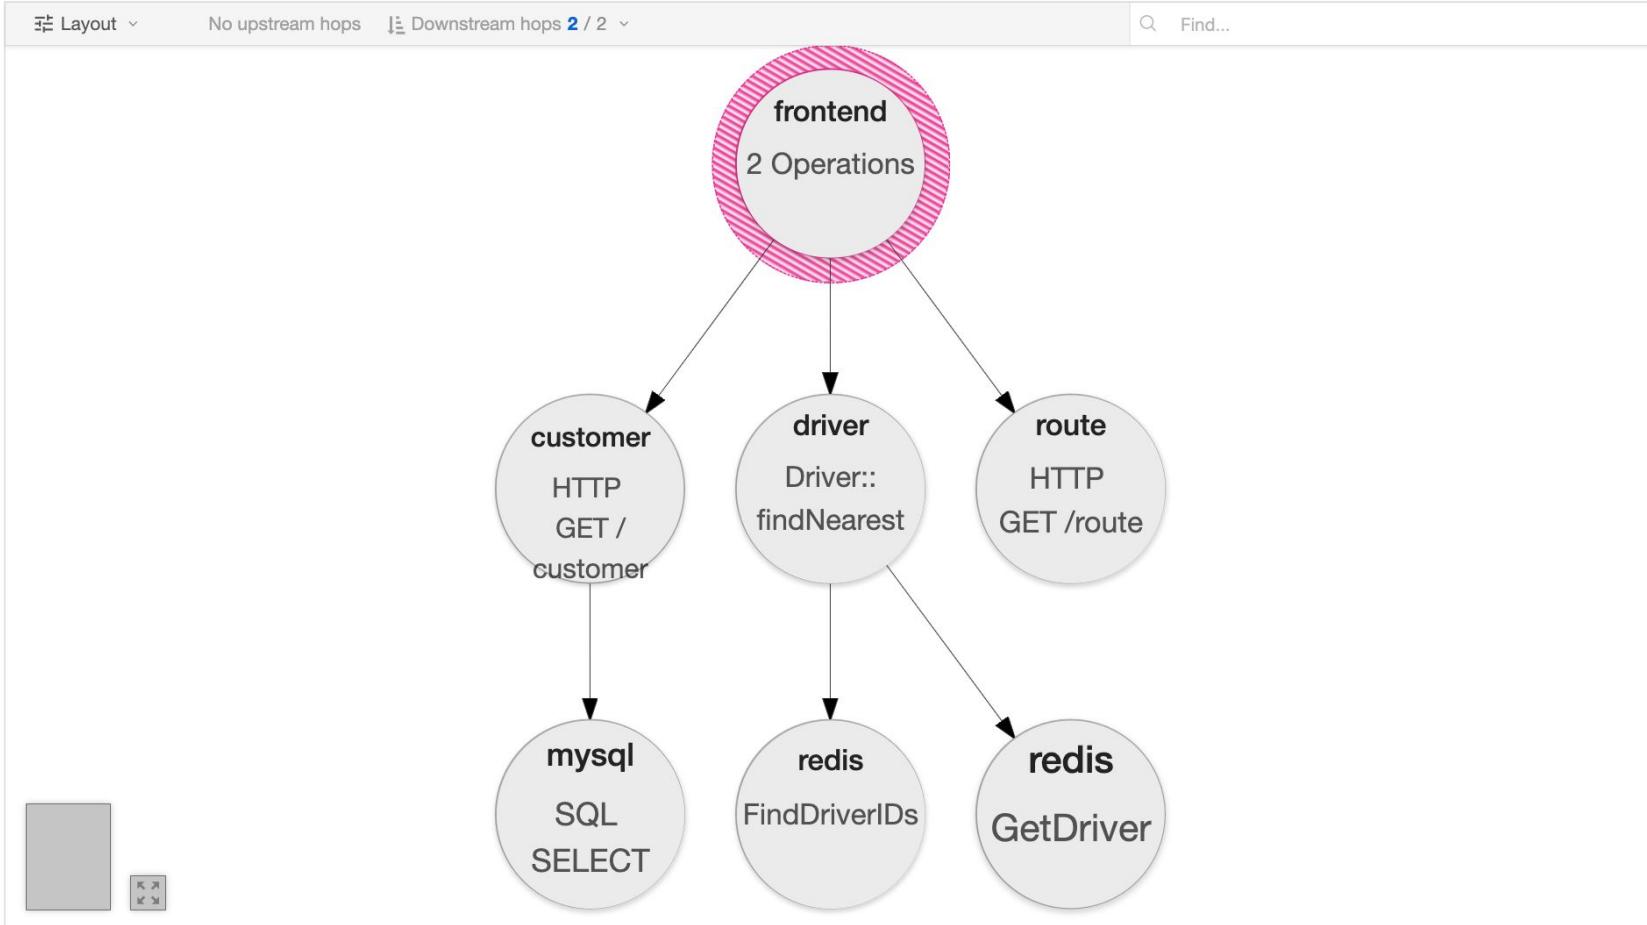
Deep Dependency Graph

Compare traces by selecting result items

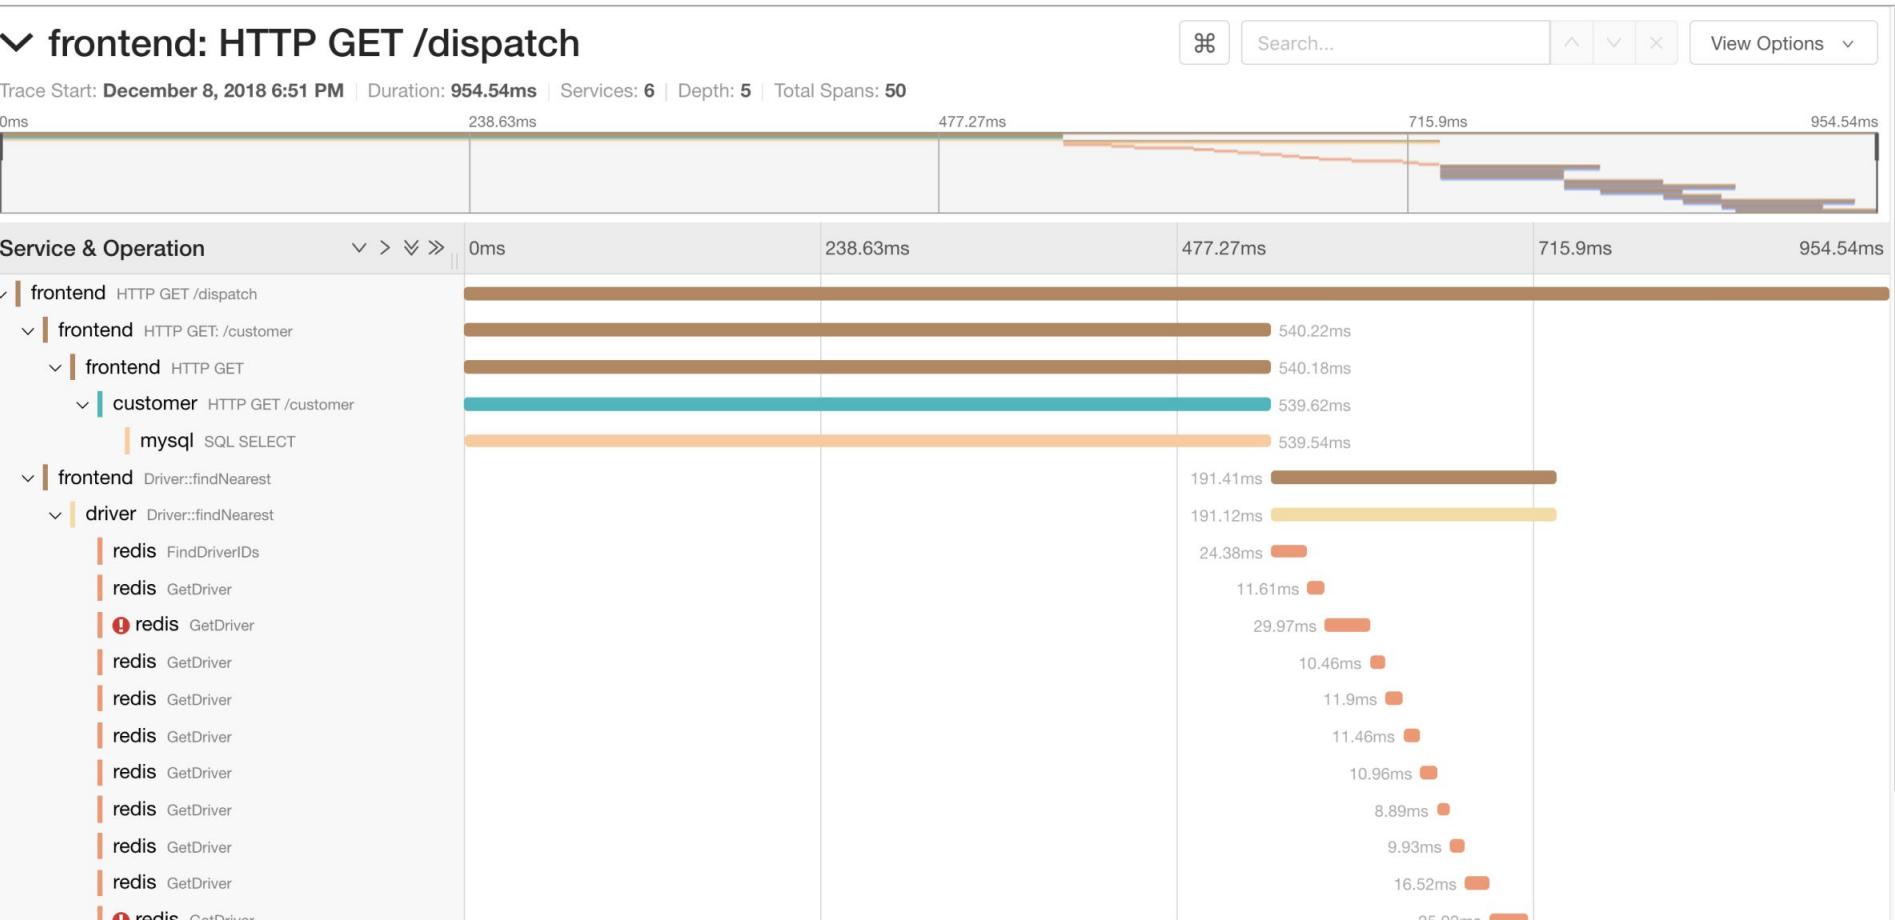


<input type="checkbox"/> frontend: HTTP GET /dispatch 3688087	1.04s
51 Spans 3 Errors customer (1) driver (1) frontend (24) mysql (1) redis (14) route (10)	Today 5:39:56 pm 5 minutes ago
<input type="checkbox"/> frontend: HTTP GET /dispatch 73e6e77	853.78ms
50 Spans 2 Errors customer (1) driver (1) frontend (24) mysql (1) redis (13) route (10)	Today 5:39:56 pm 5 minutes ago
<input type="checkbox"/> frontend: HTTP GET /dispatch d84845f	702.29ms
51 Spans 3 Errors customer (1) driver (1) frontend (24) mysql (1) redis (14) route (10)	Today 5:39:56 pm 5 minutes ago

Transitive Service Graphs



Trace timeline



Trace timeline – Parent → Child → Grandchild

✓ frontend: HTTP GET /dispatch



Search...



View Options

Trace Start: December 8, 2018 6:51 PM | Duration: 954.54ms | Services: 6 | Depth: 5 | Total Spans: 50



Service & Operation	0ms	238.63ms	477.27ms	715.9ms	954.54ms
✓ frontend HTTP GET /dispatch					
✓ frontend HTTP GET /customer				540.22ms	
✓ frontend HTTP GET				540.18ms	
✓ customer HTTP GET /customer			539.62ms		
mysql SQL SELECT			539.54ms		
✓ frontend Driver::findNearest				191.41ms	
✓ driver Driver::findNearest				191.12ms	
redis FindDriverIDs				24.38ms	
redis GetDriver				11.61ms	
redis GetDriver				29.97ms	
redis GetDriver				10.46ms	
redis GetDriver				11.9ms	
redis GetDriver				11.46ms	
redis GetDriver				10.96ms	
redis GetDriver				8.89ms	
redis GetDriver				9.93ms	
redis GetDriver				16.52ms	
redis GetDriver				25.92ms	

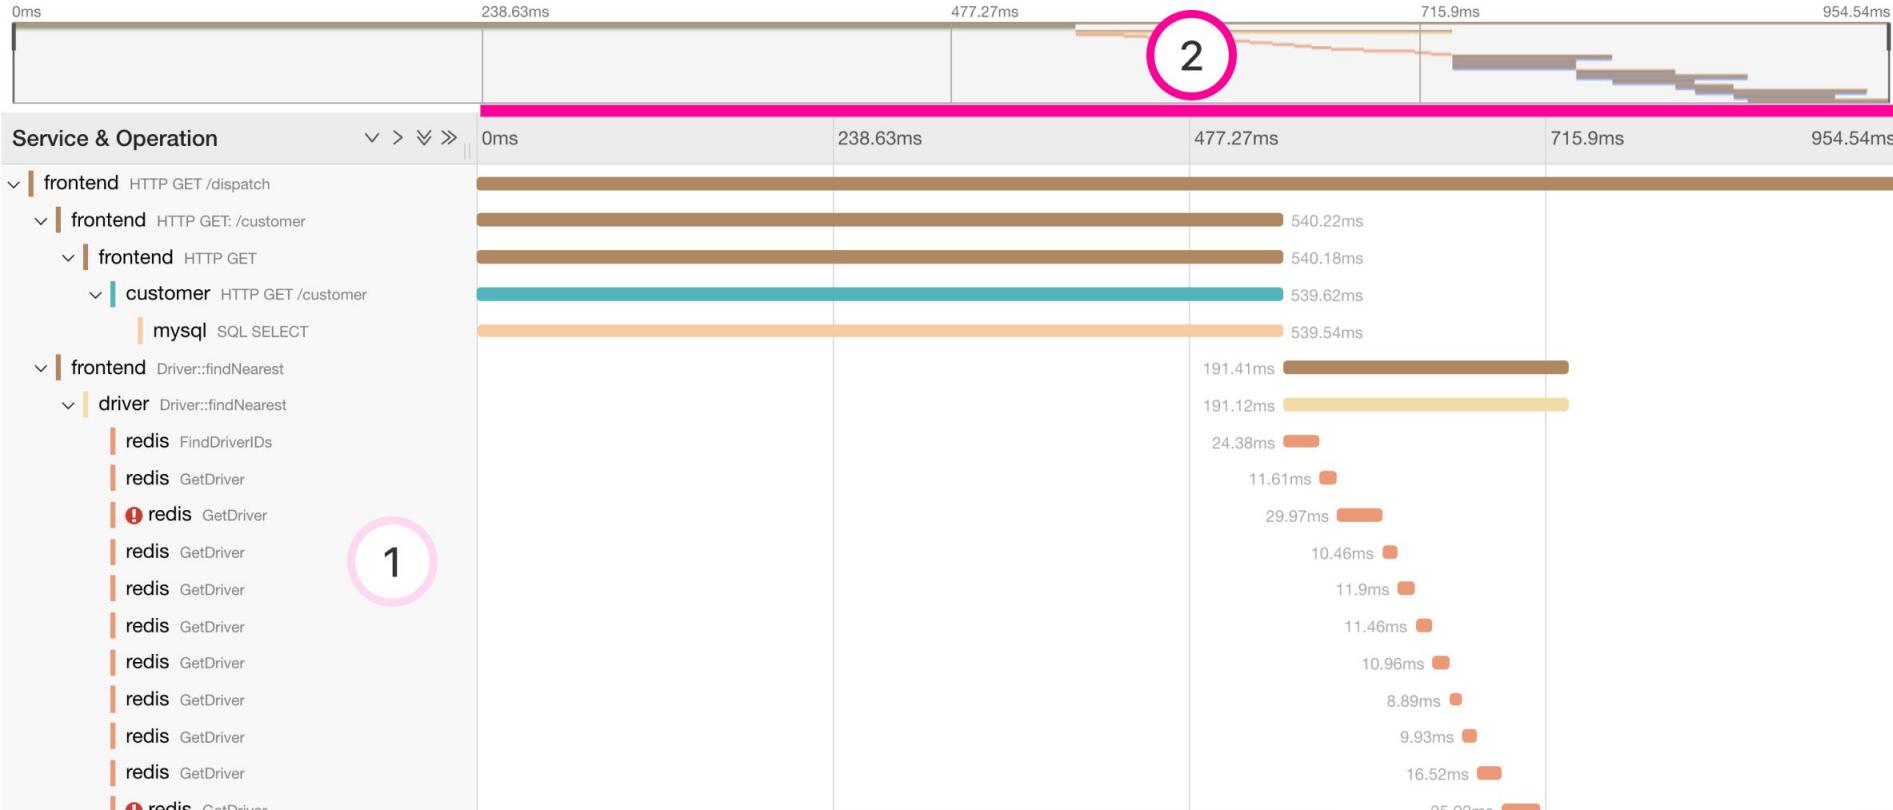
1

Trace timeline – Time + Mini-map

frontend: HTTP GET /dispatch

⌘ Search... ⌘ ⌘ ⌘ View Options ⌘

Trace Start: December 8, 2018 6:51 PM | Duration: 954.54ms | Services: 6 | Depth: 5 | Total Spans: 50



Trace timeline – A blocking operation

✓ frontend: HTTP GET /dispatch

⌘ Search... ▲ ▼ × View Options ▾

Trace Start: December 8, 2018 6:51 PM | Duration: 954.54ms | Services: 6 | Depth: 5 | Total Spans: 50



Service & Operation	0ms	238.63ms	477.27ms	715.9ms	954.54ms
✓ frontend HTTP GET /dispatch					
└ frontend HTTP GET :/customer				540.22ms	
└ frontend HTTP GET				540.18ms	
└ customer HTTP GET /customer			539.62ms		
└ mysql SQL SELECT			539.54ms		
└ frontend Driver::findNearest				191.41ms	
└ driver Driver::findNearest				191.12ms	
└ redis FindDriverIDs				24.38ms	
└ redis GetDriver				11.61ms	
└ redis GetDriver				29.97ms	
└ redis GetDriver				10.46ms	
└ redis GetDriver				11.9ms	
└ redis GetDriver				11.46ms	
└ redis GetDriver				10.96ms	
└ redis GetDriver				8.89ms	
└ redis GetDriver				9.93ms	
└ redis GetDriver				16.52ms	
└ redis GetDriver				25.92ms	

1

3

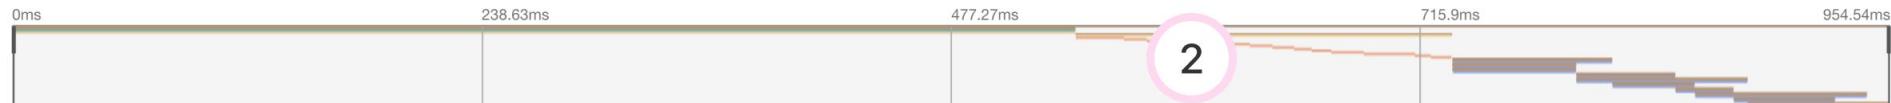
2

Trace timeline – Sequential operations

✓ frontend: HTTP GET /dispatch

⌘ Search... ⌂ ⌄ ⌅ ⌆ View Options ⌄

Trace Start: December 8, 2018 6:51 PM | Duration: 954.54ms | Services: 6 | Depth: 5 | Total Spans: 50



Service & Operation	0ms	238.63ms	477.27ms	715.9ms	954.54ms
✓ frontend HTTP GET /dispatch					
✓ frontend HTTP GET :/customer				540.22ms	
✓ frontend HTTP GET				540.18ms	
✓ customer HTTP GET /customer			539.62ms		
mysql SQL SELECT			539.54ms		
✓ frontend Driver::findNearest				191.41ms	
✓ driver Driver::findNearest				191.12ms	
redis FindDriverIDs				24.38ms	
redis GetDriver				11.61ms	
redis GetDriver				29.97ms	
redis GetDriver				10.46ms	
redis GetDriver				11.9ms	
redis GetDriver				11.46ms	
redis GetDriver				10.96ms	
redis GetDriver				8.89ms	
redis GetDriver				9.93ms	
redis GetDriver				16.52ms	
redis GetDriver				25.92ms	
redis GetDriver					

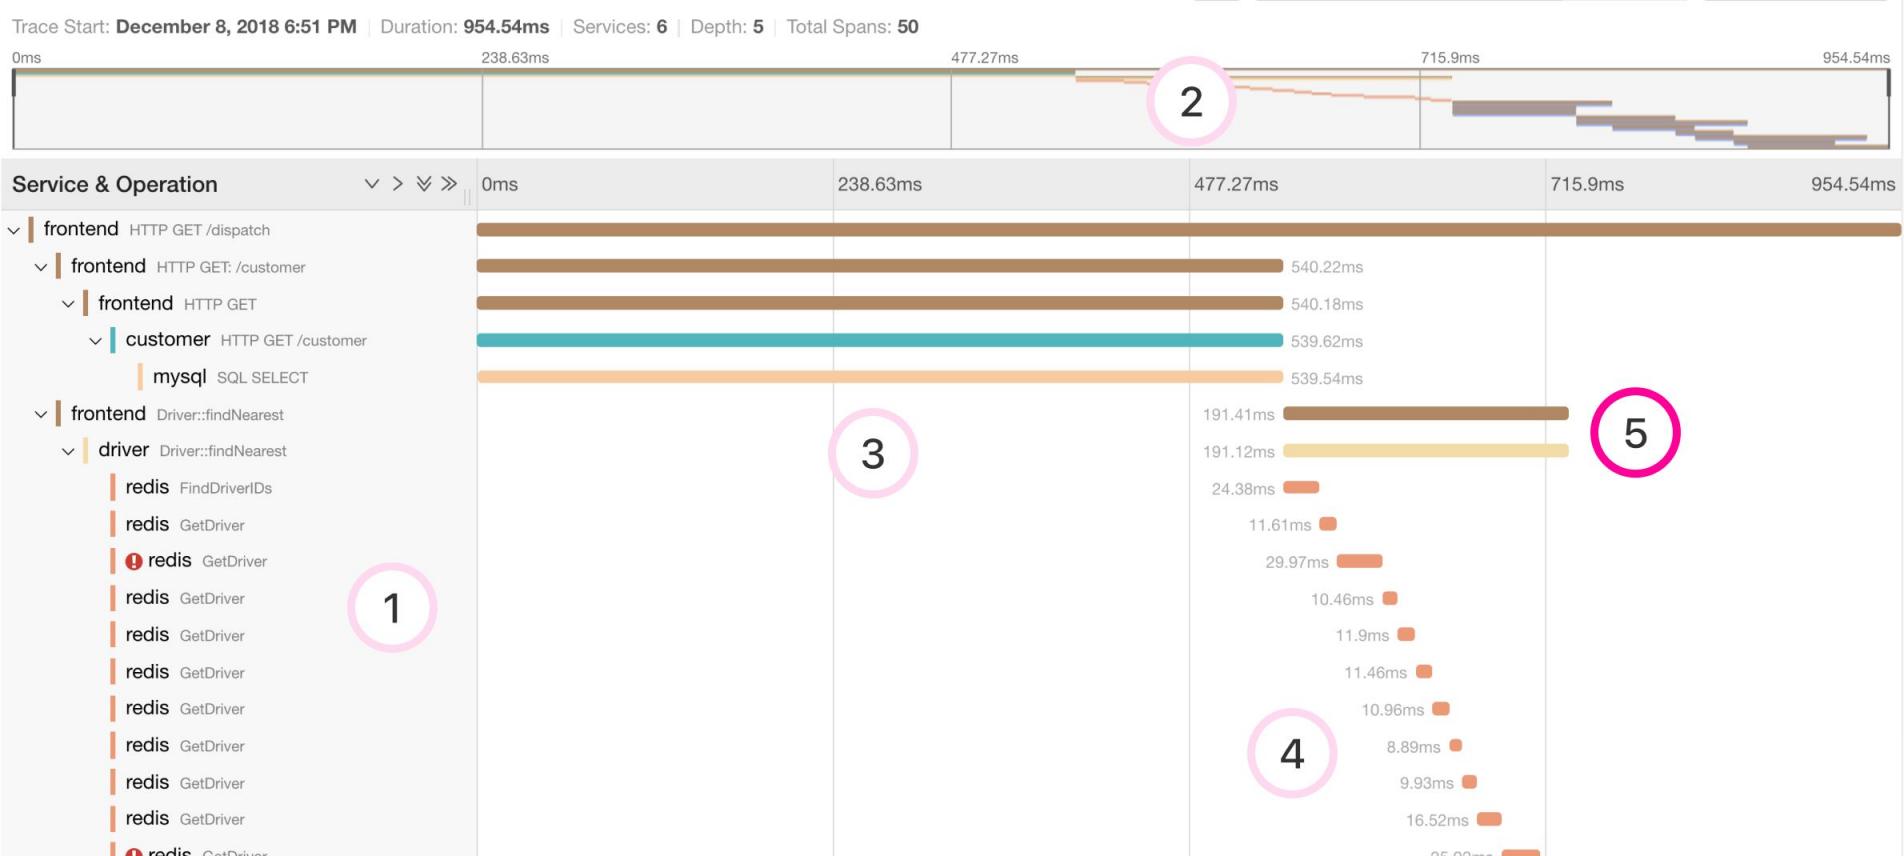
1

3

4

Trace timeline – Parents encompass descendants (generally)

✓ frontend: HTTP GET /dispatch



Span details

> frontend: HTTP GET /dispatch

Service & Operation v > v > 0ms 238.63ms 477.27ms 715.9ms 954.54ms

Search... View Options

frontend HTTP GET /dispatch

frontend HTTP GET: /customer

frontend HTTP GET

customer HTTP GET /customer

mysql SQL SELECT

SQL SELECT Service: mysql | Duration: 539.54ms | Start Time: 0.67ms

Tags

span.kind	"client"
peer.service	"mysql"
sql.query	"SELECT * FROM customer WHERE customer_id=392"
request	"3878-3"

Process: client-uuid = 55627059ae2defbd | hostname = joef-C02TX0LYHTDG | ip = 192.168.1.5 | jaeger.version = Go-2.15.0

Logs (2)

0.68ms: event = Waiting for lock behind 2 transactions blockers = [3878-1 3878-2]

282.29ms: event = Acquired lock with 0 transactions waiting behind

Log timestamps are relative to the start time of the full trace.

SpanID: 7aecd811f9df684

frontend Driver::findNearest

driver Driver::findNearest

Span details – Database query

> frontend: HTTP GET /dispatch

Service & Operation 0ms 238.63ms 477.27ms 715.9ms 954.54ms

frontend HTTP GET /dispatch
 frontend HTTP GET: /customer
 frontend HTTP GET
 customer HTTP GET /customer
mysql SQL SELECT

SQL SELECT

Service: mysql | Duration: 539.54ms | Start Time: 0.67ms

Tags

- span.kind "client"
- peer.service "mysql"
- sql.query "SELECT * FROM customer WHERE customer_id=392"
- request "3878-3"

1

Process: client-uuid = 55627059ae2defbd | hostname = joef-C02TX0LYHTDG | ip = 192.168.1.5 | jaeger.version = Go-2.15.0

Logs (2)

- 0.68ms: event = Waiting for lock behind 2 transactions blockers = [3878-1 3878-2]
- 282.29ms: event = Acquired lock with 0 transactions waiting behind

Log timestamps are relative to the start time of the full trace.

SpanID: 7aecd811f9df684

frontend Driver::findNearest
 driver Driver::findNearest

191.41ms
191.12ms

Span details – Lock contention

> frontend: HTTP GET /dispatch

Service & Operation v > v > 0ms 238.63ms 477.27ms 715.9ms 954.54ms

- frontend HTTP GET /dispatch
- frontend HTTP GET: /customer
 - frontend HTTP GET
 - customer HTTP GET /customer
- mysql SQL SELECT

SQL SELECT Service: mysql | Duration: 539.54ms | Start Time: 0.67ms

Tags

- span.kind "client"
- peer.service "mysql"
- sql.query "SELECT * FROM customer WHERE customer_id=392"
- request "3878-3"

1

Process: client-uuid = 55627059ae2defbd | hostname = joef-C02TX0LYHTDG | ip = 192.168.1.5 | jaeger.version = Go-2.15.0

Logs (2)

0.68ms: event = Waiting for lock behind 2 transactions blockers = [3878-1 3878-2]

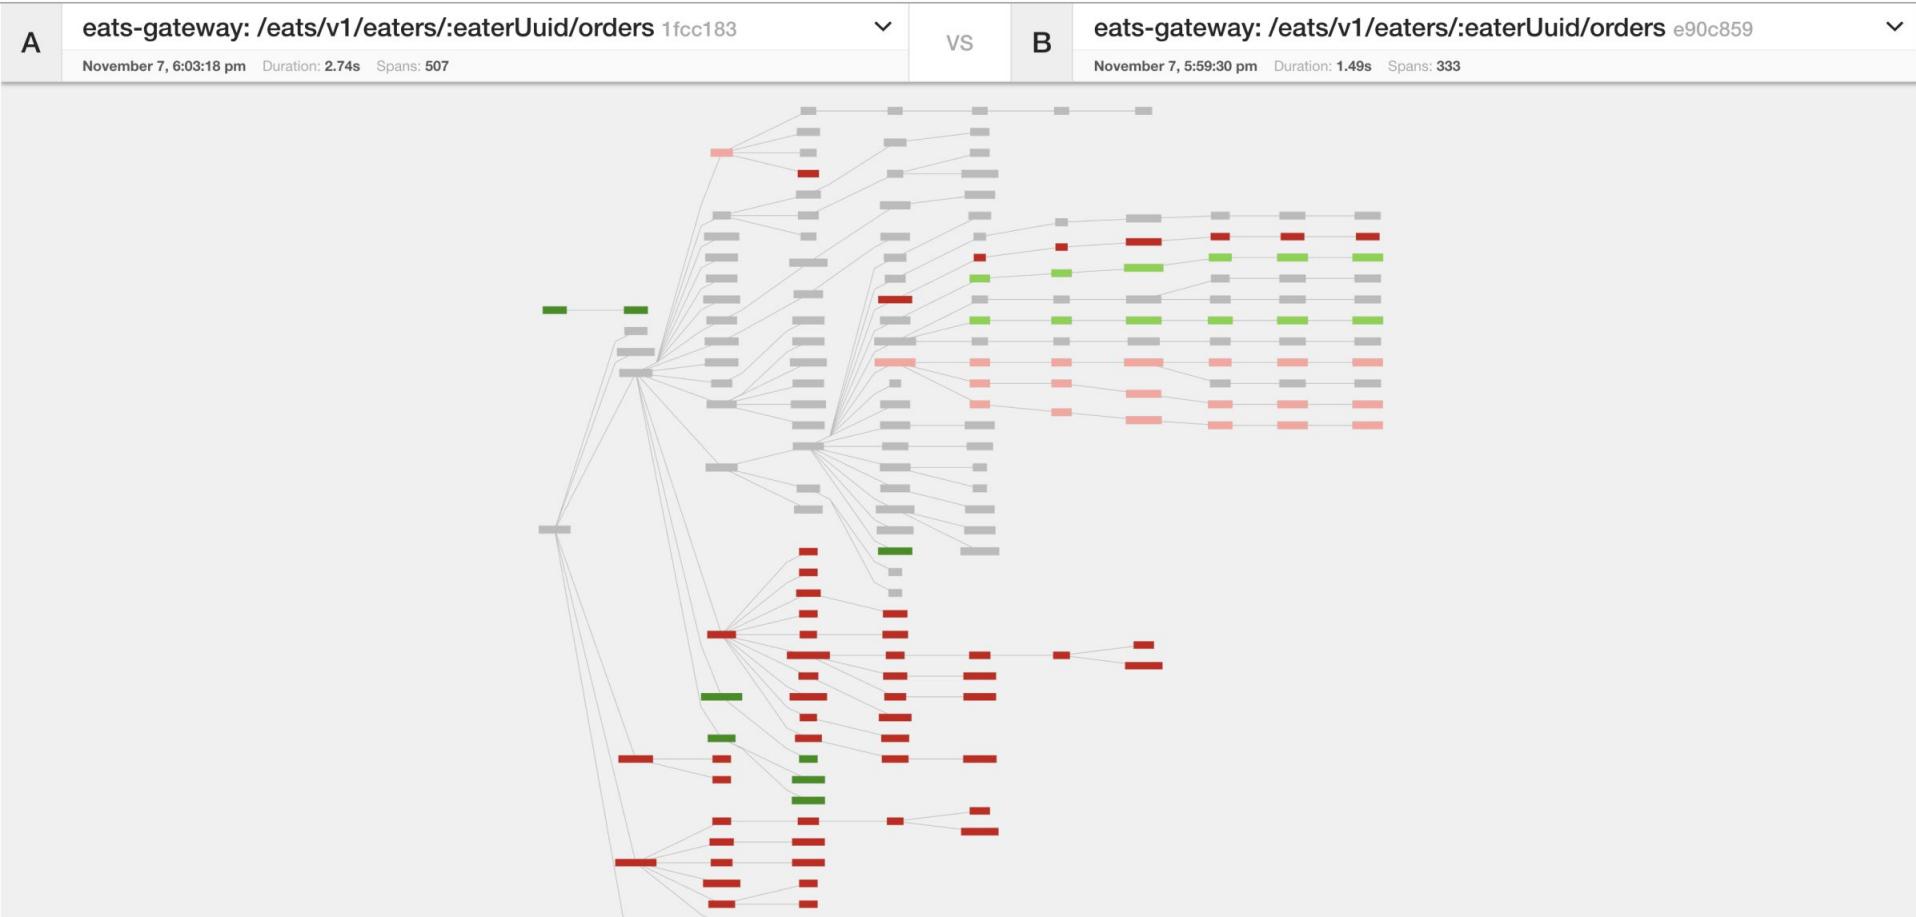
282.29ms: event = Acquired lock with 0 transactions waiting behind

Log timestamps are relative to the start time of the full trace.

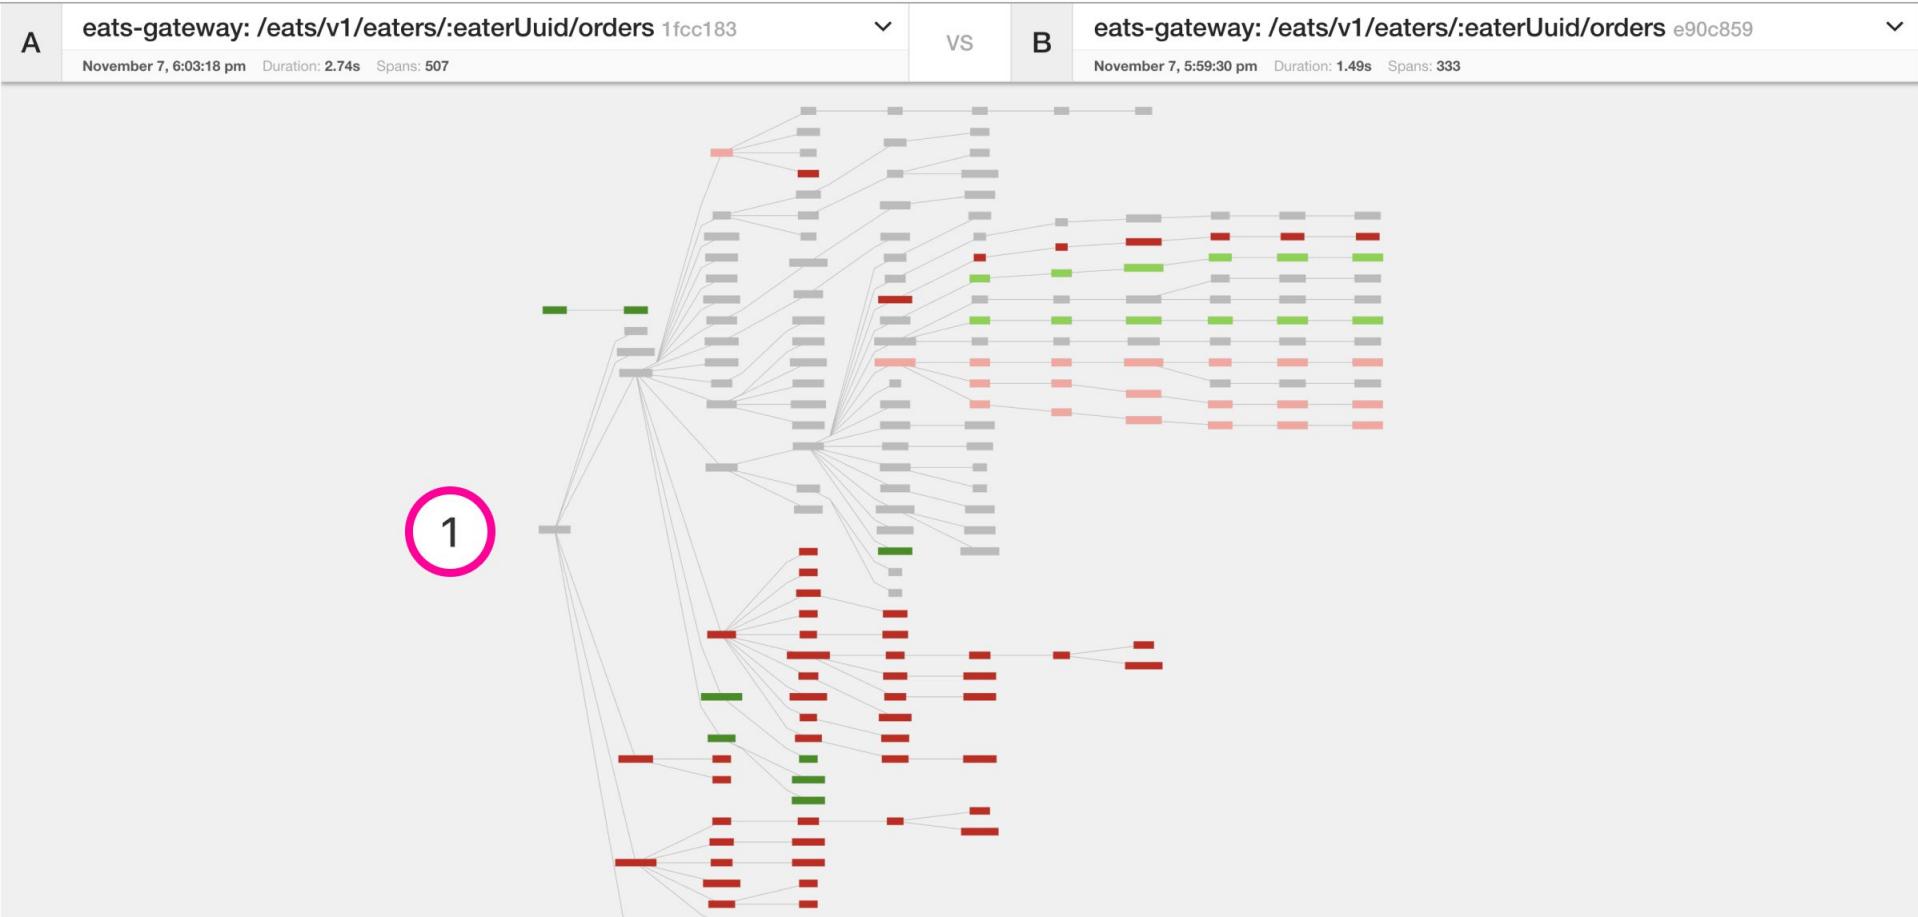
SpanID: 7aecd811f9df684

- frontend Driver::findNearest
- driver Driver::findNearest

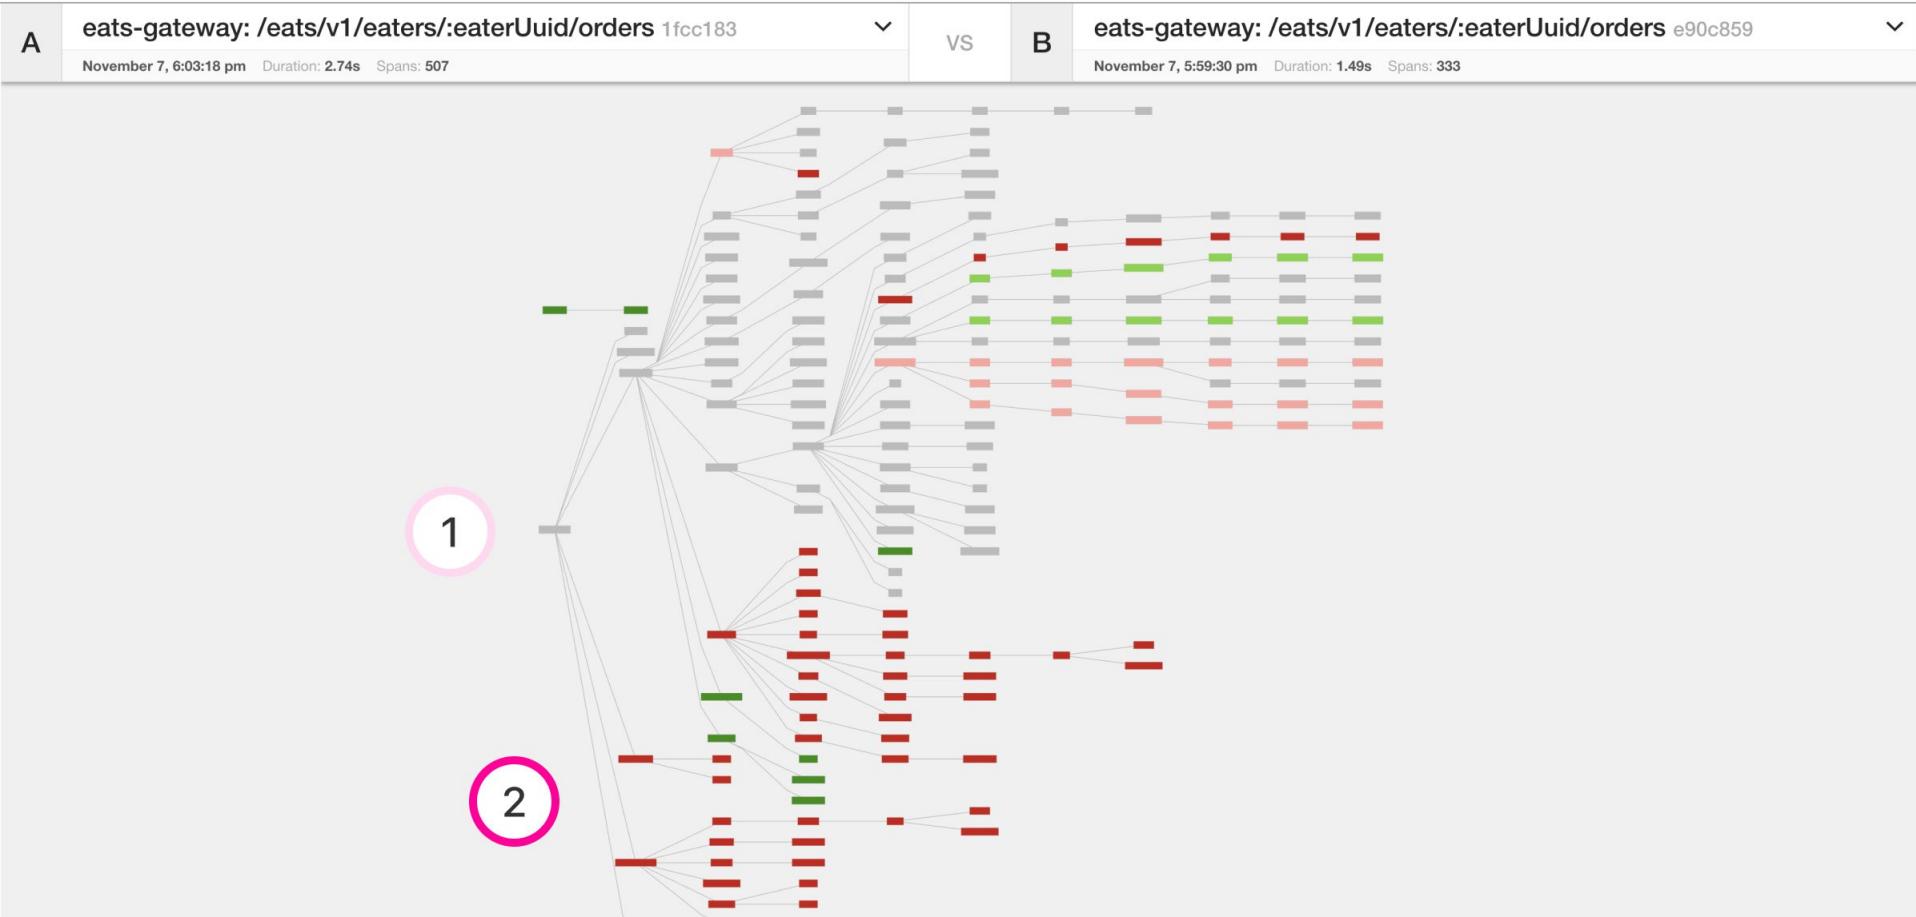
Comparing trace structures – Unified diff



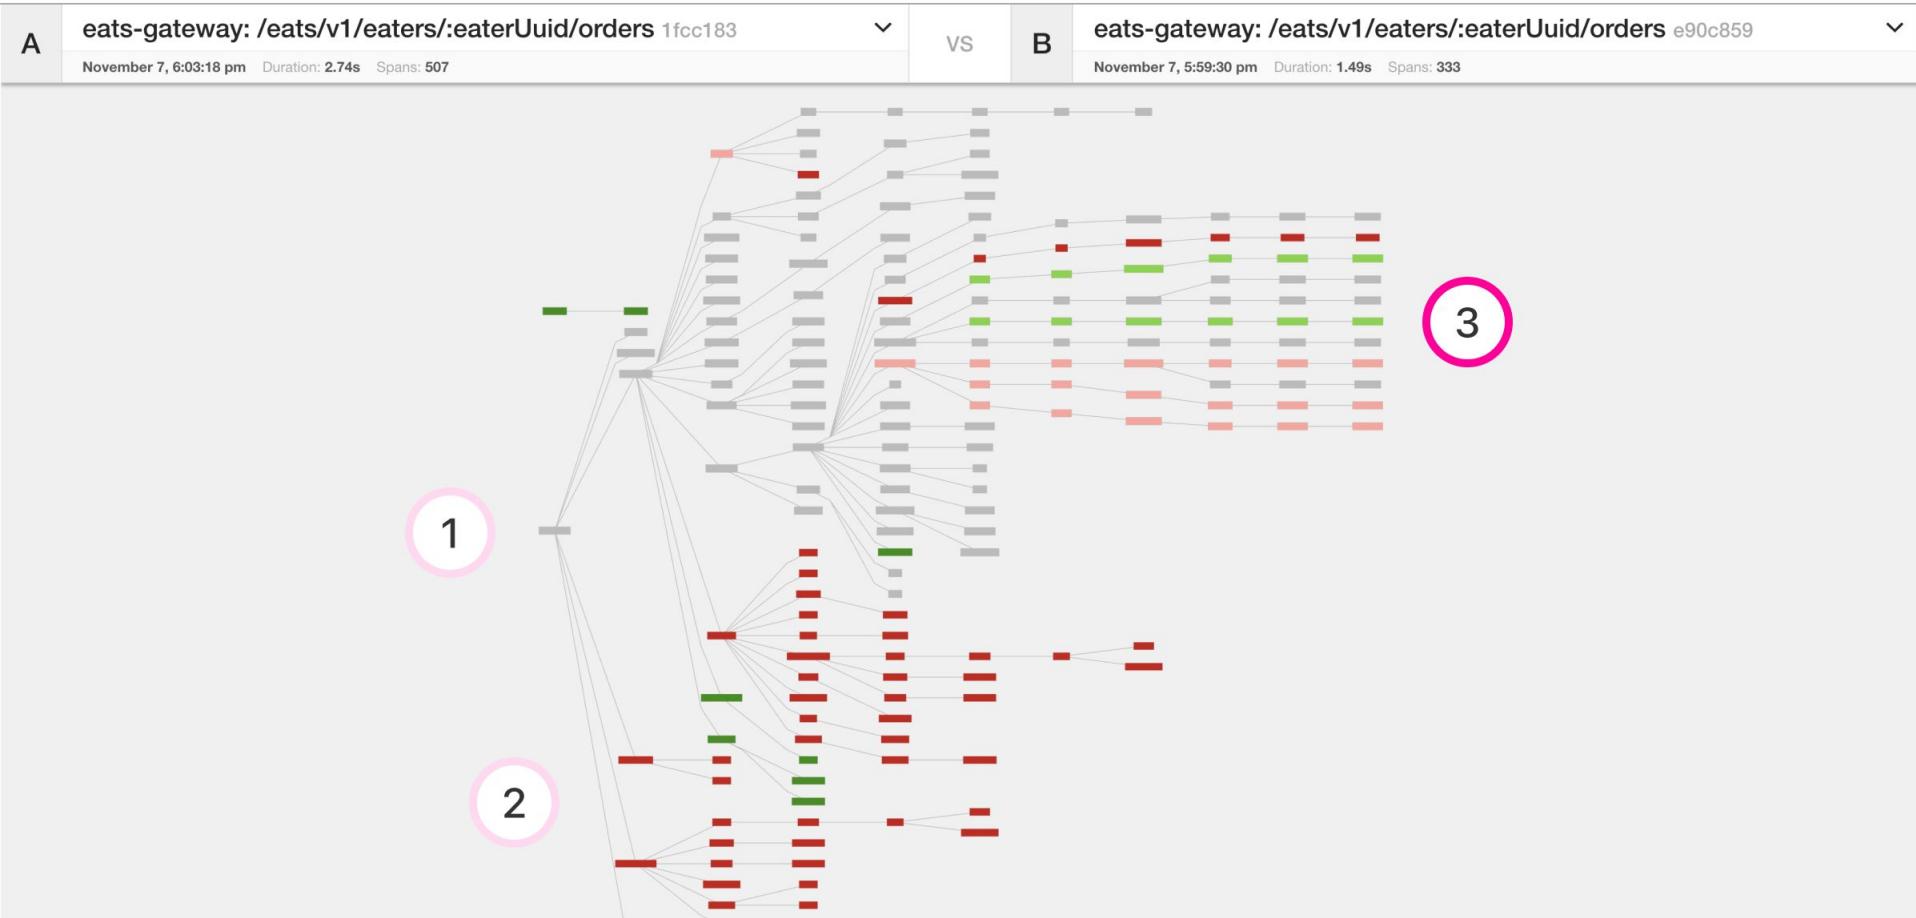
Comparing trace structures – Shared structure



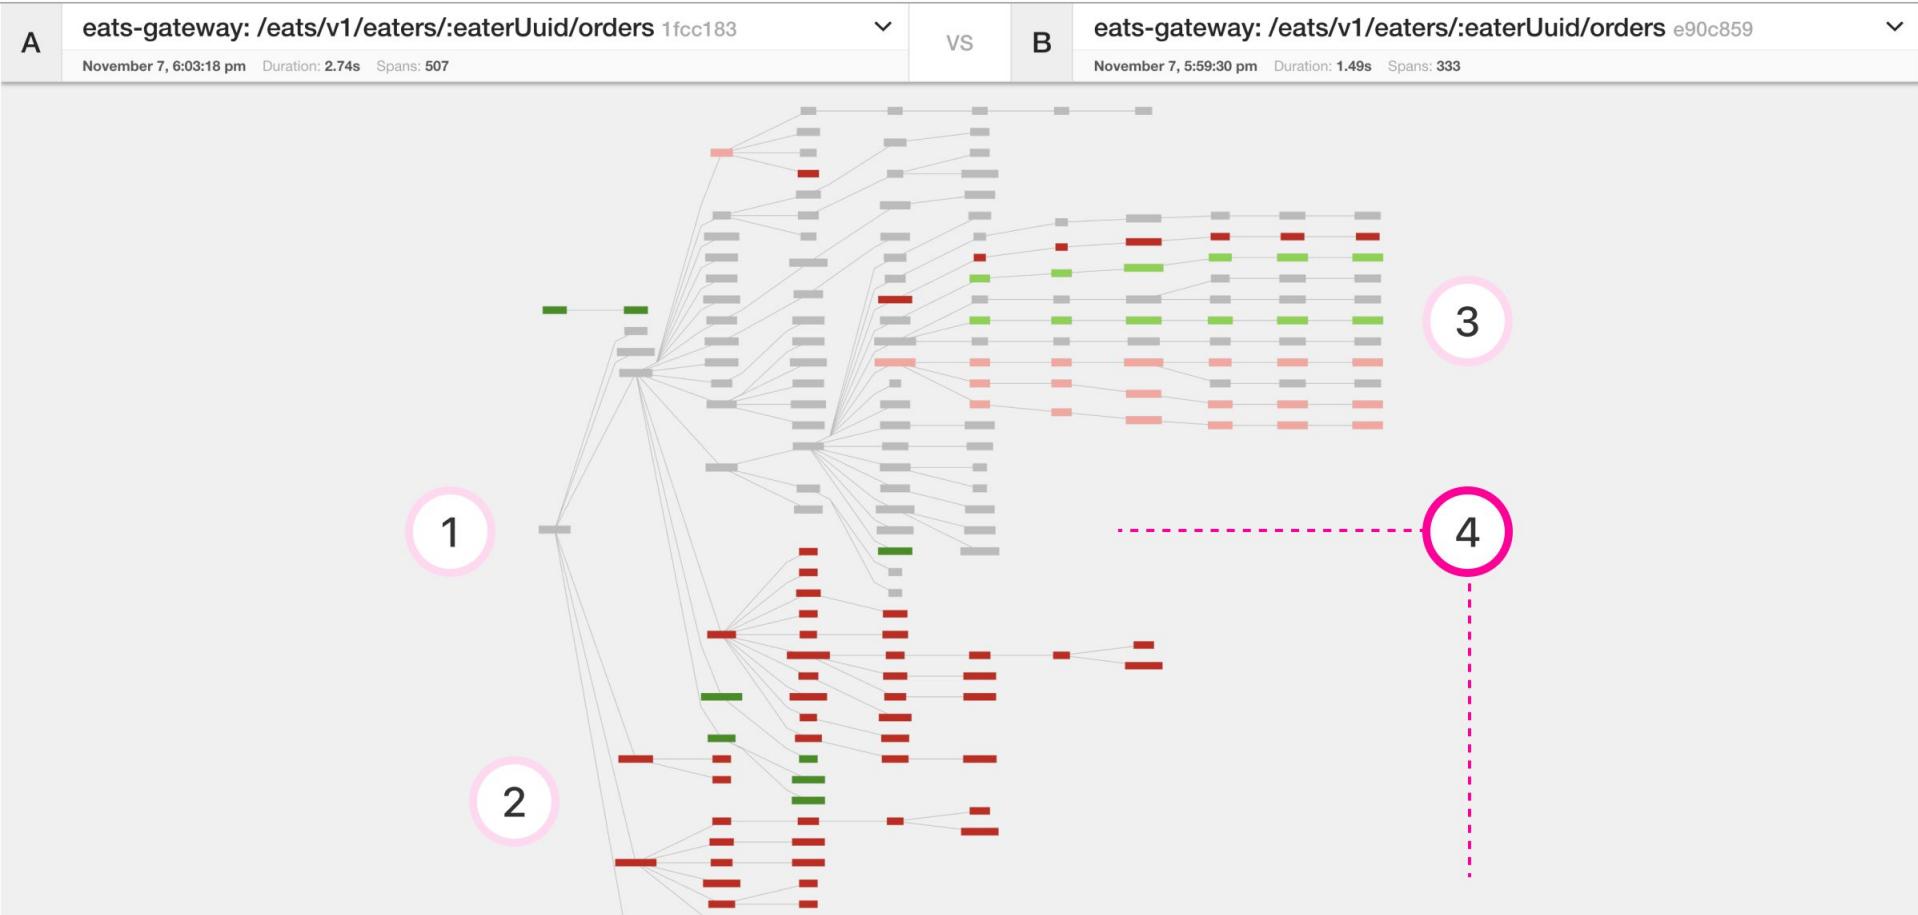
Comparing trace structures – Absent in one or the traces



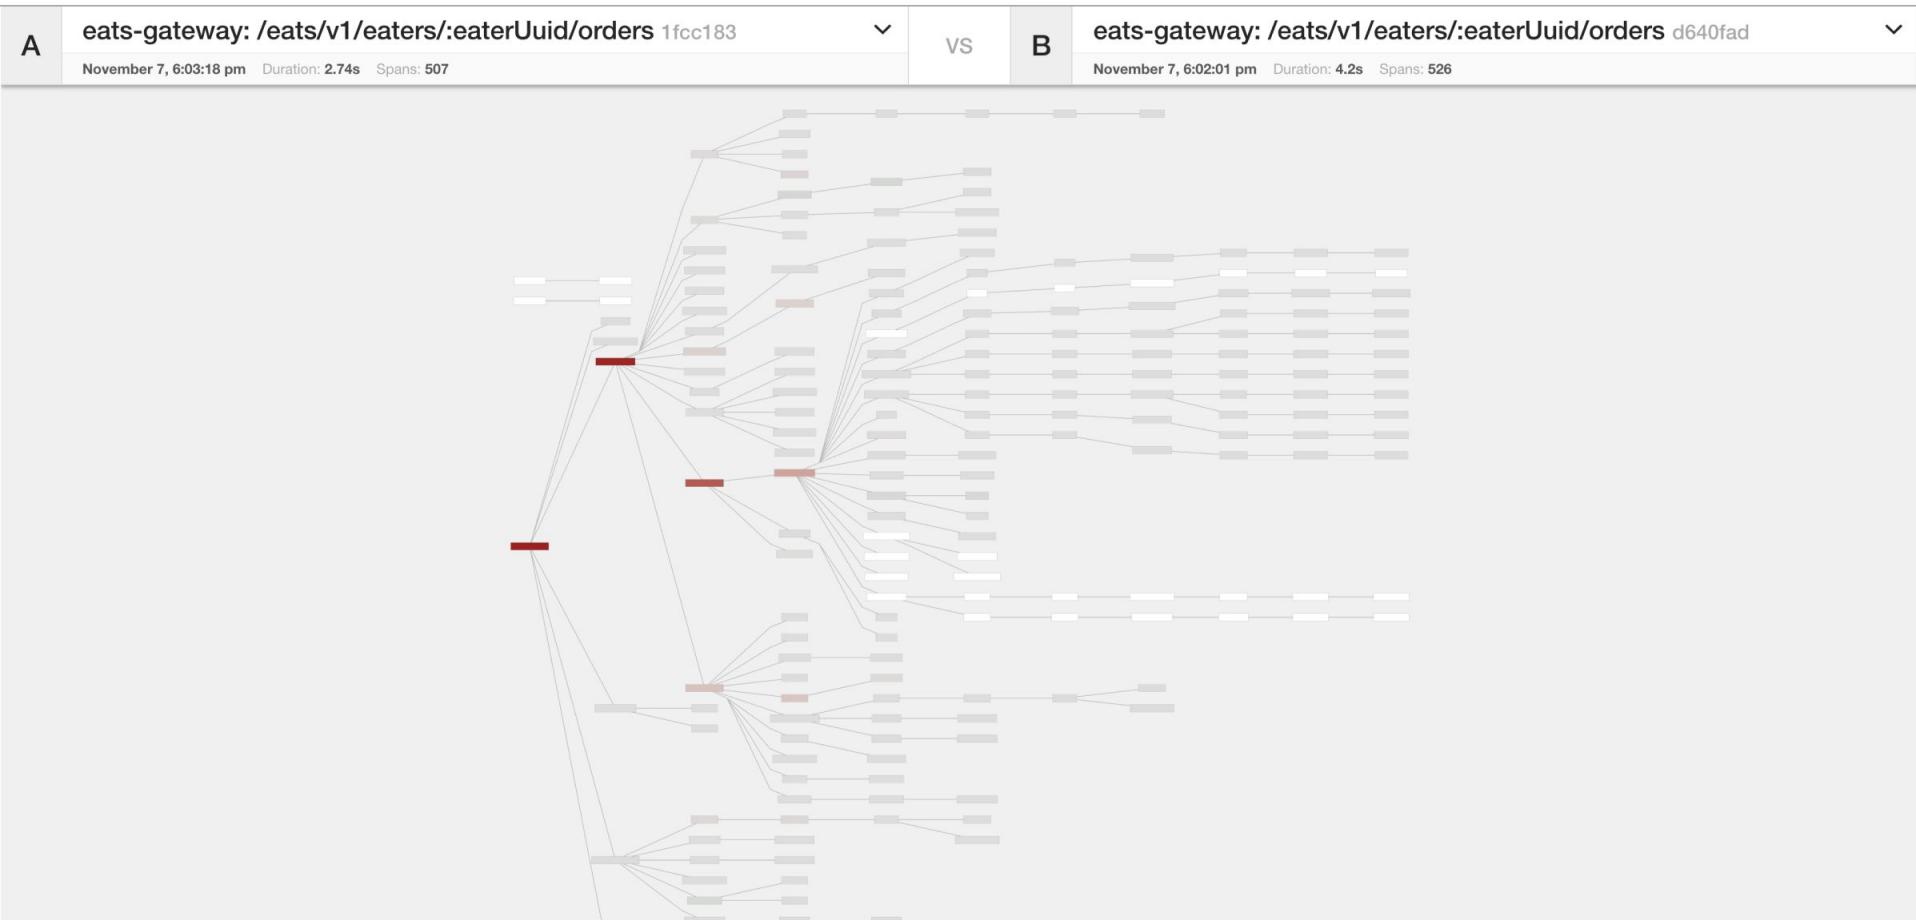
Comparing trace structures – More or less within a node



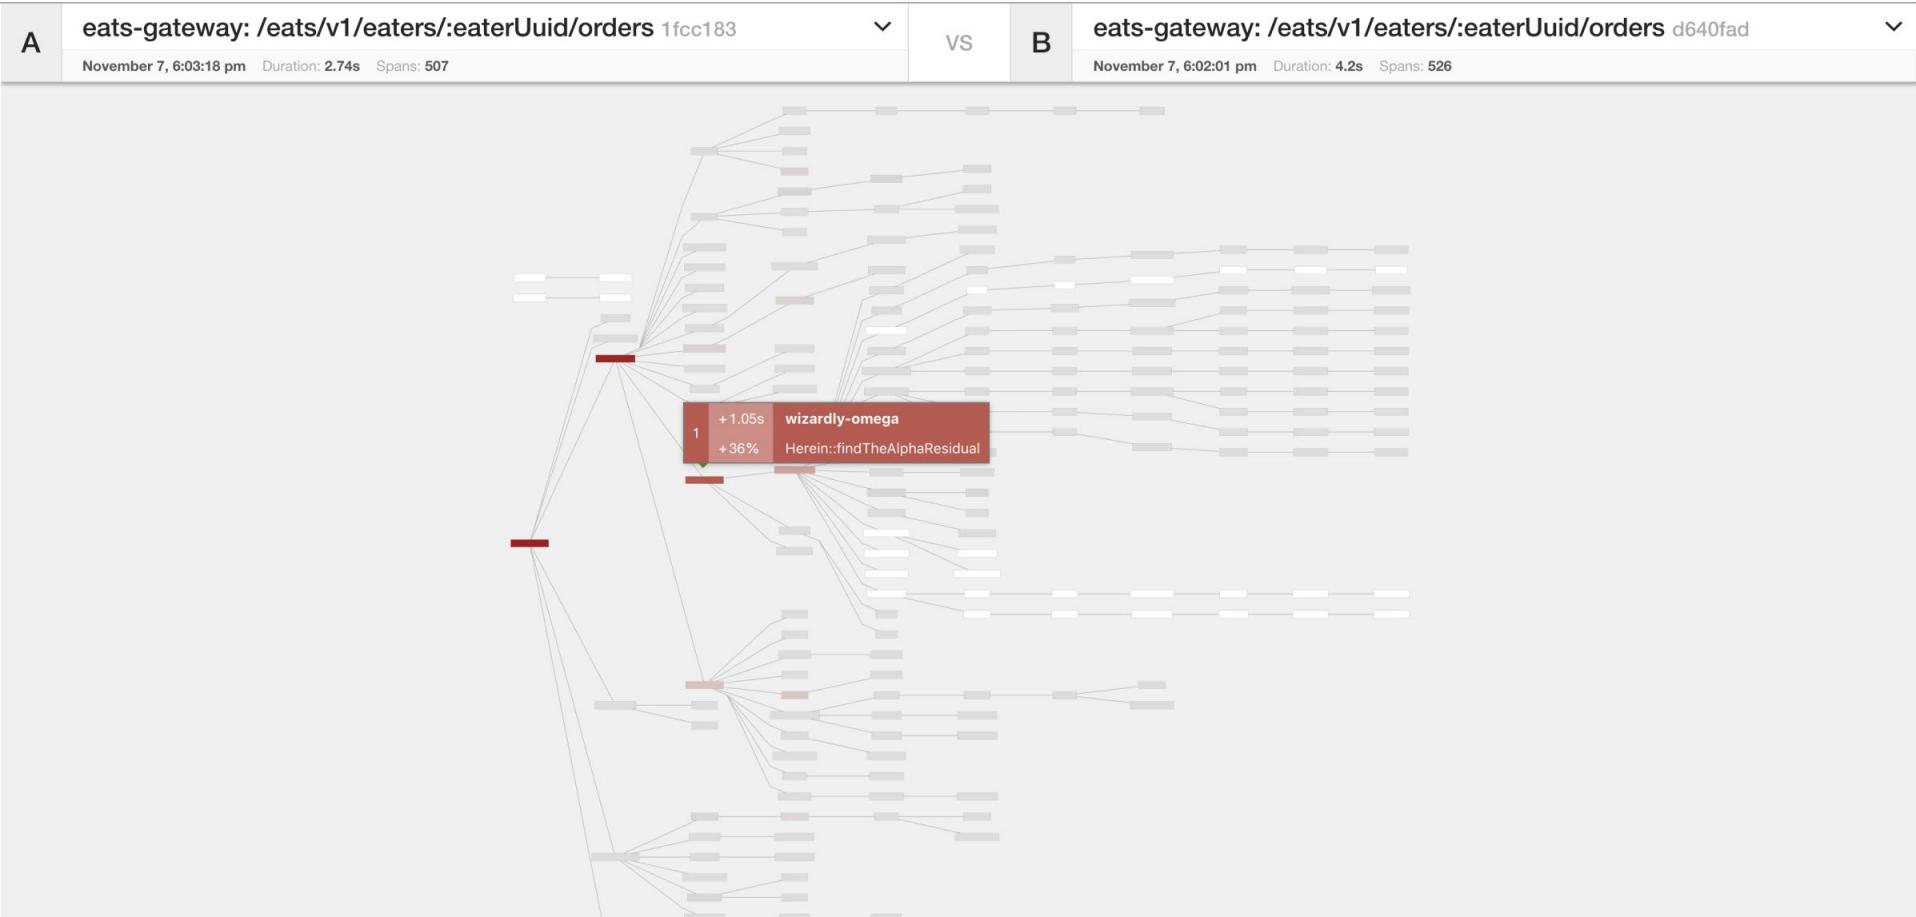
Comparing trace structures – Substantial divergence



Comparing span durations



Comparing span durations





CLOUD NATIVE
COMPUTING FOUNDATION

Jaeger Architecture

Jaeger, a Distributed Tracing Platform



Instrumentation not included

Jaeger project does not provide instrumentation!

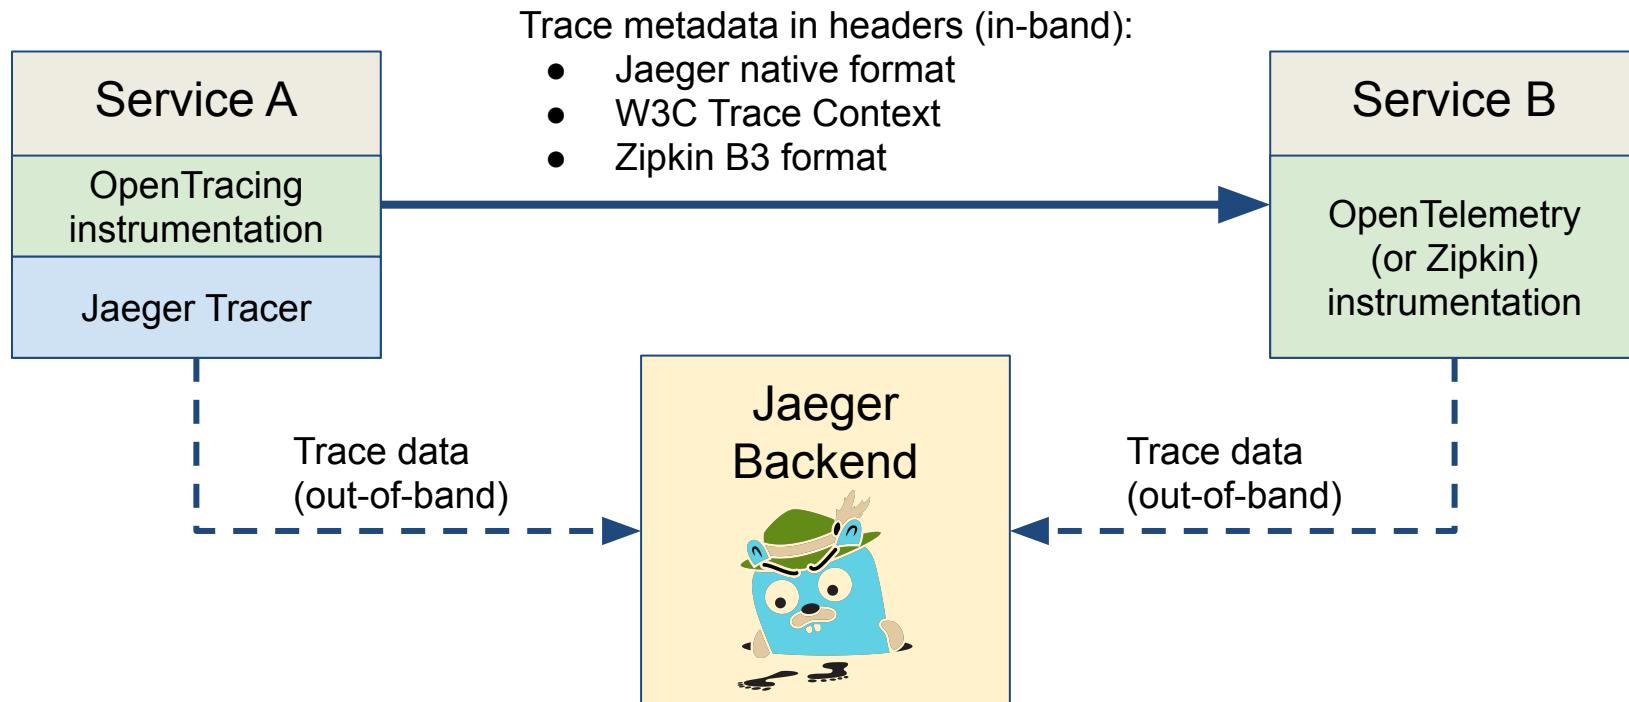
Use OpenTracing or OpenTelemetry.

Jaeger - /'yāgər/, noun: hunter

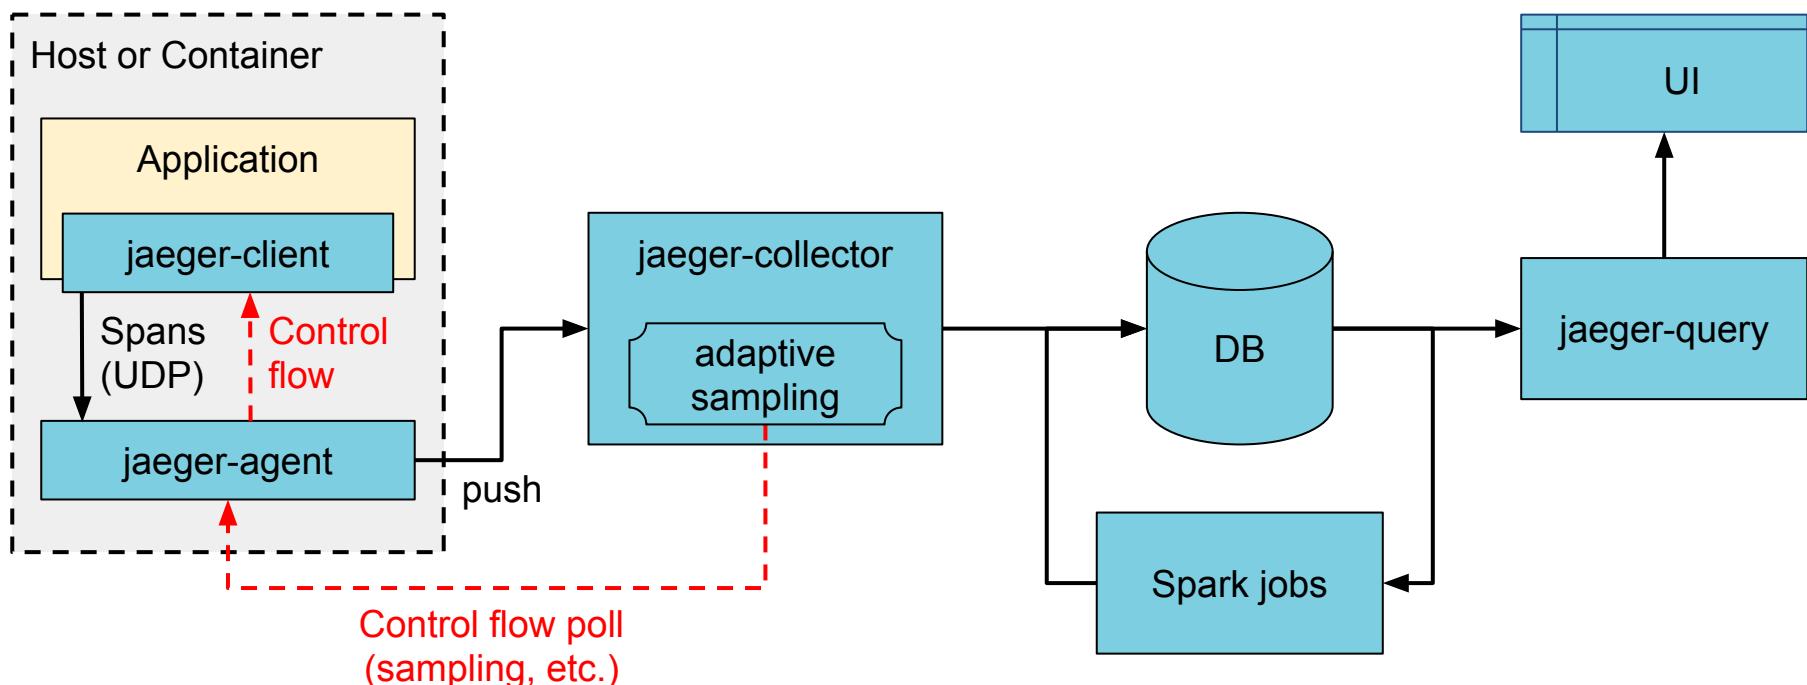
- Inspired by Google's Dapper and OpenZipkin
- Created at Uber in August 2015 ([blog](#))
- Open sourced in April 2017
- Joined CNCF in Sep 2017 (as incubating)
- Graduated to top-level CNCF project Oct 2019 ([CNCF announcement](#))



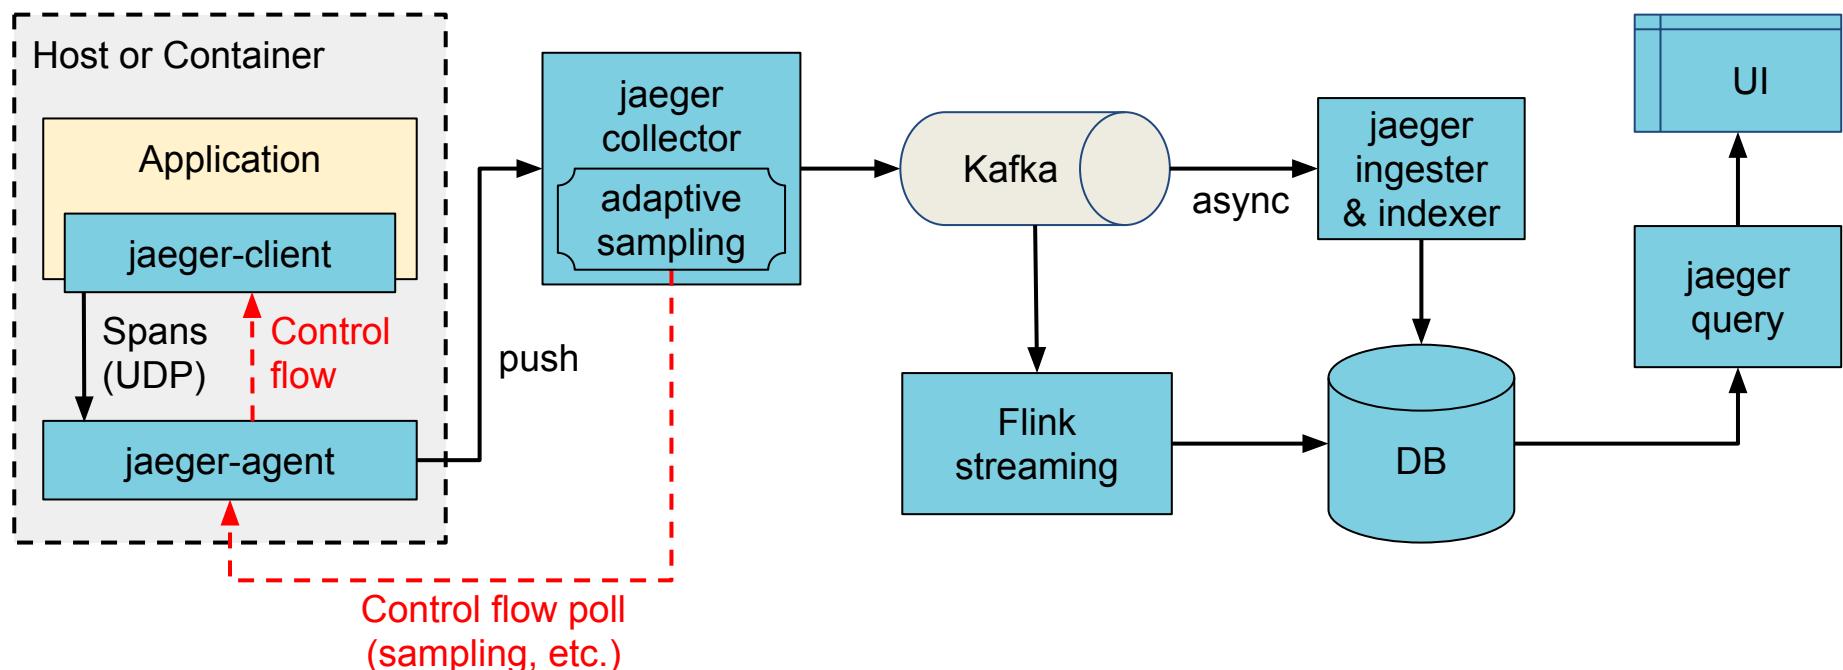
Jaeger and your system



Architecture 2017: Push



Architecture now: Push+Async+Streaming



Technology Stack

- Go backend
- Pluggable storage
 - Cassandra, Elasticsearch, badger, memory
- React/Javascript frontend
- OpenTracing Instrumentation libraries
- Integration with Kafka, Apache Flink



Go



Java™
POWERED

python
powered





Jaeger

And Sampling



Sampling

Sampling is the selection of a subset (a statistical sample) of individuals from within a (statistical) population to estimate characteristics of the whole population.

reason about application performance

traces

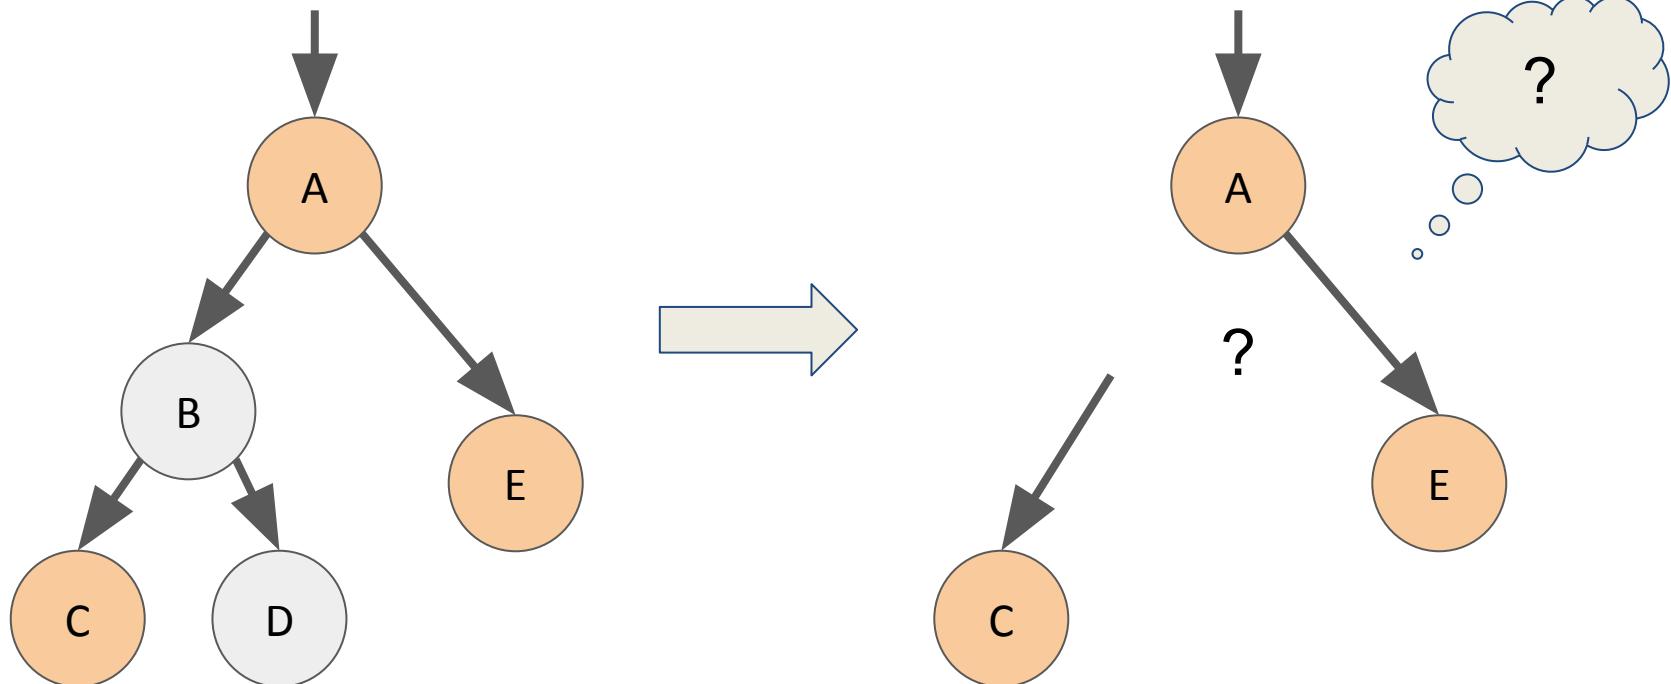
all possible traces

Why do we sample

1. Saving everything incurs large storage costs
 - 2 KB / span on 10k QPS server \Rightarrow 20 MB/s
 - $x100$ instances \Rightarrow 2 GB/s \approx 170 PB/day (for one service!)
2. Performance overhead from instrumentation
 - 10k QPS server \Rightarrow 100 μ s / req budget
 - Trace instrumentation: 5 μ s \Rightarrow 5% overhead
3. Trace data is very repetitive



Goal: Consistent (all or nothing) Sampling



Sampling techniques

- Head-based sampling
 - Most popular in the industry
- Tail-based sampling
 - Gaining popularity recently

Head-based (upfront) sampling

Sampling decision is made at the start of the trace and propagated in the trace context.

- ✓ Minimal perf overhead when trace is not sampled
- ✓ Easy to implement, supported by Jaeger SDKs
- ✗ Can easily miss rare anomalies/outliers
 - Prob. of catching p99 latency with 1% sampling rate $\Rightarrow 1/10,000$
- ✗ Cannot “sample on errors”

Head-based sampling in Jaeger

- SDKs can be configured with different samplers (always on / off, probabilistic, rate limiting, etc.)
 - ✓ Easy to implement
 - ✗ Spread-out configuration in the hands of developers
- SDKs default to “remote” sampler that allows centralized configuration (polled from collectors)

Jaeger sampling configuration

```
"default_strategy": {  
    "type": "probabilistic",  
    "param": 0.5,  
    "operation_strategies": [  
        {  
            "operation": "/health",  
            "type": "probabilistic",  
            "param": 0.0  
        },  
        {  
            "operation": "/metrics",  
            "type": "probabilistic",  
            "param": 0.0  
        }  
    ]  
}
```

Applies to all other services

```
"service_strategies": [  
    {  
        "service": "foo",  
        "type": "probabilistic",  
        "param": 0.8,  
        "operation_strategies": [  
            {  
                "operation": "bar",  
                "type": "probabilistic",  
                "param": 0.2  
            }  
        ]  
    }  
]
```

Custom strategy per service

Overrides for specific endpoints

Overrides for specific endpoints

Tail-based (post-trace) sampling

Sampling decision is made at the end of the trace:

- ✓ Can be much more intelligent, based on observed latency, errors, unusual call patterns & graph shapes, etc.
- ✓ Can catch anomalies
- ✓ Can perform aggregations before sampling
- ✗ Requires temporary storage of all traces
- ✗ Applications incur performance overhead even for traces that may be later discarded

Tail-based sampling in Jaeger

- ✓ Supported in jaeger-opentelemetry-collector
- ✓ Configurable sampling rules: latency, certain tags
- ✗ Single-node mode only, multi-node sharded solution will be available in the future

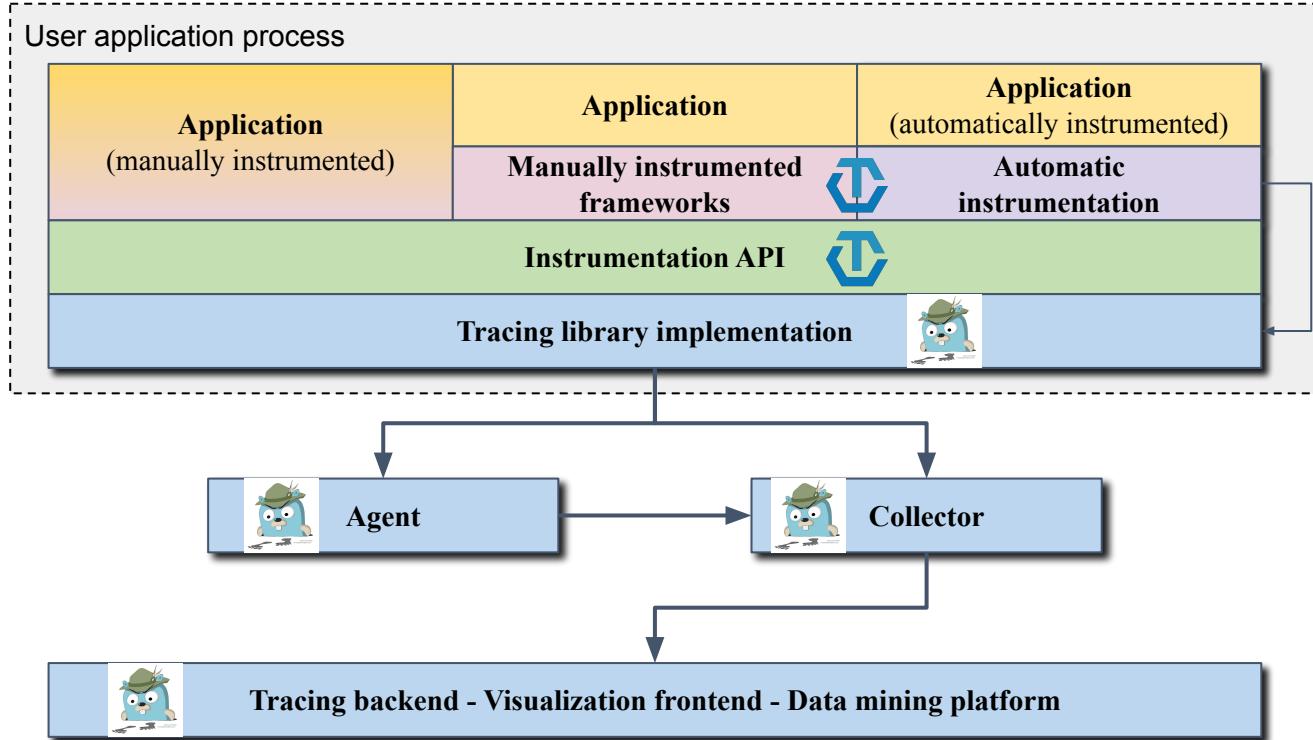


Jaeger

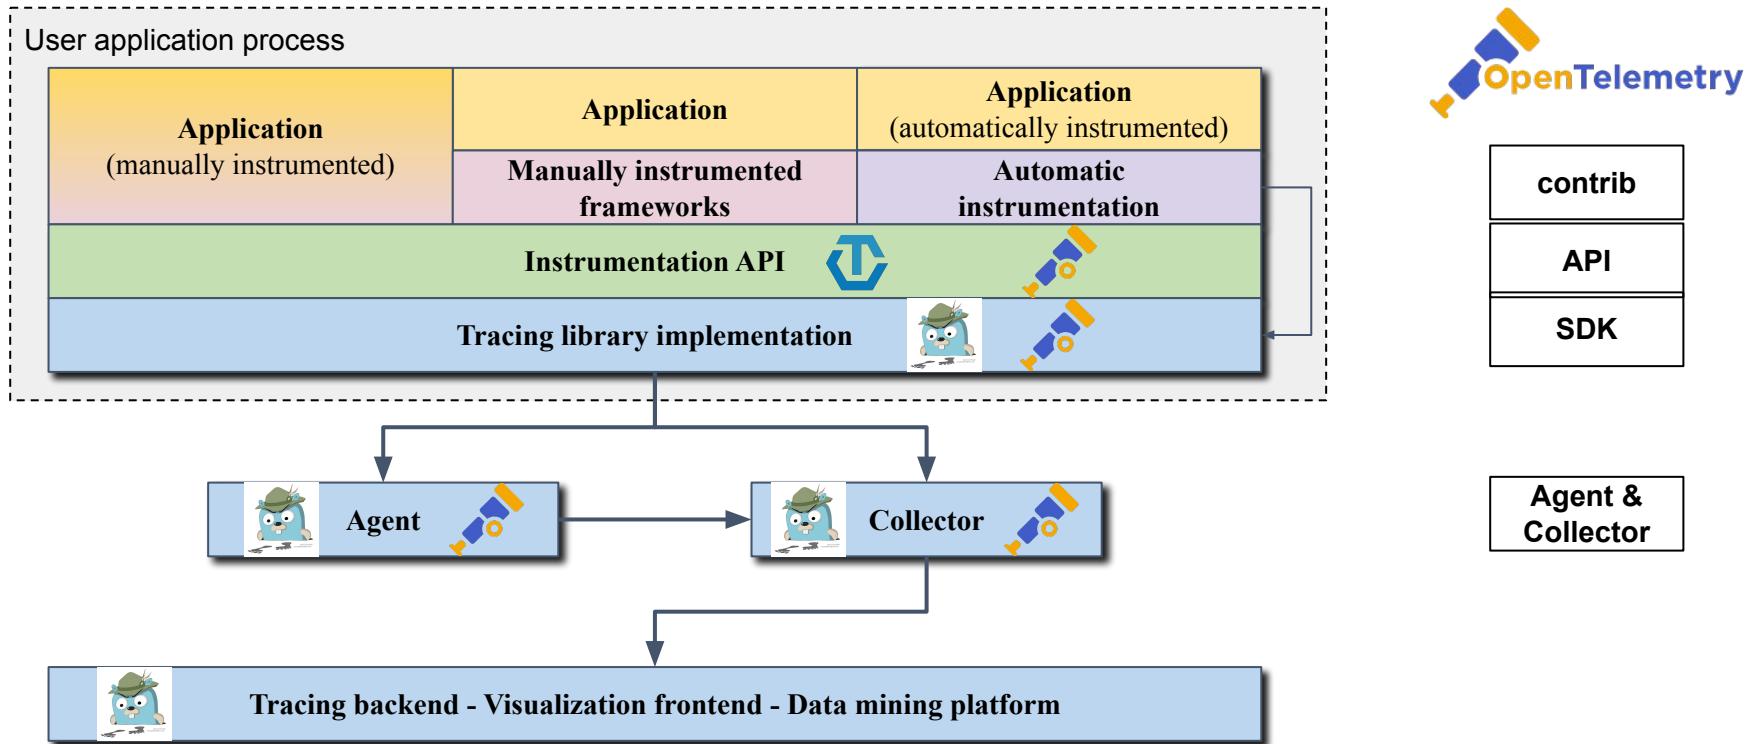
And OpenTelemetry



Jaeger with OpenTracing



Jaeger with OpenTelemetry



Jaeger components on OpenTelemetry

- OpenTelemetry Collector is written in Go
- We built Jaeger-specific versions
 - Have the same capabilities as upstream OTEL
 - With Jaeger extensions, e.g. storage
- We're converting Jaeger storage implementation to OTEL data model for better compatibility



Jaeger

And OpenTelemetry SDKs



Jaeger and OpenTelemetry SDKs

- OpenTelemetry SDK support
 - Jaeger gRPC exporter
 - Jaeger propagation
- OpenTracing SHIM
 - use OTel SDK with OpenTracing instrumentations
- Jaeger client libraries support W3C Trace Context



Jaeger

And Kubernetes



Deploying Jaeger on Kubernetes

- Helm charts
- Jaeger Operator
 - allInOne and production deployment
 - auto provisioning of Kafka (Strimzi)
- Plain Kubernetes manifest files

Getting in Touch

- GitHub: <https://github.com/jaegertracing>
- Chat: <https://gitter.im/jaegertracing/>
- Mailing List - jaeger-tracing@googlegroups.com
- Blog: <https://medium.com/jaegertracing>
- Twitter: <https://twitter.com/JaegerTracing>
- Bi-Weekly Community Meetings