

Stop Writing Operators

Joe Thompson

Senior Solutions Engineer, HashiCorp



Stop Writing Operators

by default ^

Joe Thompson

Senior Solutions Engineer, HashiCorp





In IT since my first job helping out with computers in my high school in 1994

Past employers: Mesosphere, Capital One, CoreOS, Red Hat, among others

Exposed to Kubernetes in early 2015 and working with it full-time since late 2015

Currently a Solutions Engineer for  HashiCorp (we're hiring!)

Pronouns: he/him
Blood type: Caffeine-positive

Contact info:

 jthompson@hashicorp.com

 Twitter: [@caffeinepresent](https://twitter.com/caffeinepresent)

 Kubernetes Slack: [@kensey](https://kubernetes.slack.com)

[LinkedIn](#)





KubeCon



CloudNativeCon

North America 2020

Virtual

28 Jun
2020

Talk
submitted



12 Sep 2020

Devan Goodwin: When Not to
Write a Kubernetes Operator



22 Sep 2020

The New Stack: Kubernetes: When to Use, and
When to Avoid, the Operator Pattern

29 Sep
2020

Talk
accepted!





KubeCon



CloudNativeCon

North America 2020

Virtual

- Should this operator be written at all?
- Should I write it?
- What functions should my operator perform?
- What API permissions will my operator need to perform those functions?
- What framework should I write my operator in?



Virtual

Are you serious?

- Pretty serious
- Not like, we-can't-be-friends-unless serious though

Are you saying operators are bad?

- No! Operators are good!
...in the proper places

What an operator is and isn't



Virtual

2016 -- CoreOS: Introducing Operators: Putting Operational Knowledge into Software
"...application-specific operational knowledge encoded into software
that leverages the powerful Kubernetes abstractions to run and
manage the application correctly..."

Meant for managing stateful apps

- All operators are controllers but NOT all controllers are operators!
- A controller that does nothing stateful is just a controller (and that's OK)

Deciding if you want an operator

Kubernetes



Your app



KubeCon



CloudNativeCon

— North America 2020 —

Virtual

No statefulness

- or -

Kubernetes has all the necessary state handling



KubeCon



CloudNativeCon

North America 2020

Virtual

Creating an operator is over-abstracting the problem

You now have an extra component susceptible to failure



KubeCon



CloudNativeCon

North America 2020

Virtual

Isn't an operator just another microservice?

- No: microservices remain abstracted from one another. Operators inherently don't.



Eric Chiang
@erchiang

Things learned last week: Kubernetes operators are a huge pain for security reviews. You basically have to reverse engineer the app to figure out what holes it'll open in your cluster.

1:41 PM · Jan 13, 2020 · Twitter Web App

27 Retweets 2 Quote Tweets 172 Likes



Eric Chiang @erchiang · Jan 13
Replying to @erchiang

There are a handful of apps that need operators (e.g. databases). For the rest, please, please provide an option to deploy from static manifests. If a component requests permissions to create RBAC cluster roles and bindings, how can a security reviewer reasonably assess that?

3

5

30



Jay Vyas @jayunit100 · Jan 14
Replying to @erchiang

Use audit2rbac and MAC in those cases , thanks to @liggitt for creating it !



Jimmy Zelinskie @jimmyzelinskie · Jan 13
Replying to @erchiang

This is exactly why CoreOS and Red Hat have been trying to build a framework around operators. Once you have a bunch of them you realize how painful all of this is.

Operators (and CRD controllers) make your security burden heavier

Alternatives to operators



KubeCon



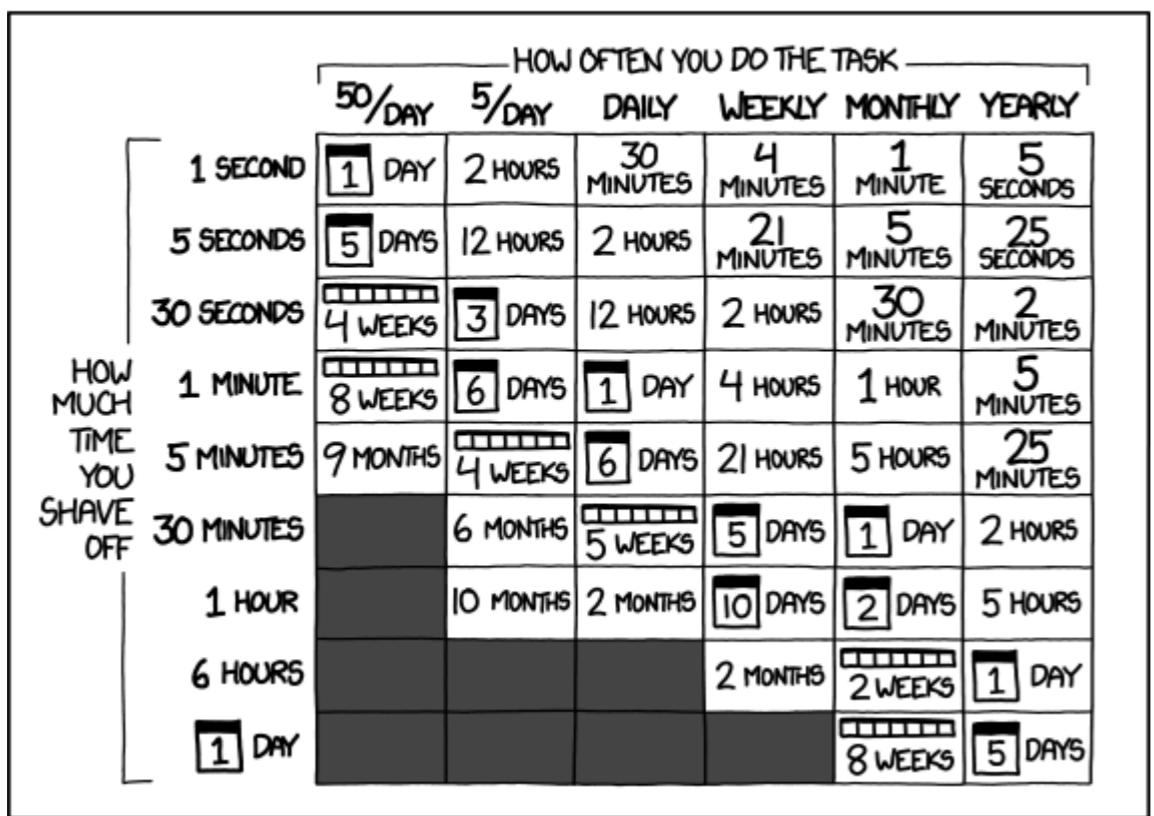
CloudNativeCon

North America 2020

Virtual

Not trying to manage state at all is easier to do safely

HOW LONG CAN YOU WORK ON MAKING A ROUTINE TASK MORE
EFFICIENT BEFORE YOU'RE SPENDING MORE TIME THAN YOU SAVE?
(ACROSS FIVE YEARS)



(source: XKCD #1205)

Do nothing

- Automation isn't worth it if you aren't saving time

[Code](#)[Issues 3.2k](#)[Pull requests 428](#)[Actions](#)[Projects](#)[Wiki](#)[Security](#)[Insights](#)[master](#)[1 branch](#)[1 tag](#)[Go to file](#)[Add file](#)[Code](#)**kelseyhightower** add style guide

3 6c073b0 on Jan 21 4 commits

[CONTRIBUTING.md](#)

add no code

3 years ago

[Dockerfile](#)

add Docker support

3 years ago

[LICENSE](#)

add no code

3 years ago

[README.md](#)

add windows support

3 years ago

[STYLE.md](#)

add style guide

9 months ago

README.md

No Code

No code is the best way to write secure and reliable applications. Write nothing; deploy nowhere.

About

The best way to write secure and reliable applications. Write nothing; deploy nowhere.

[Readme](#)[Apache-2.0 License](#)

Releases 1

[1.0.0 Latest](#)

on Feb 6, 2018

Packages

No packages published

GitOps - Operations by Pull Request

At Weaveworks, developers are responsible for operating our Weave Cloud SaaS. “GitOps” is our name for how we use developer tooling to drive operations. This post talks about GitOps, which is 90% best practices and 10% cool new stuff that we needed to build. Fair warning: this is what works for us, and dear reader, you may disagree.

Git is a part of every developer’s toolkit. Using the practices outlined in this post, our developers operate Kubernetes via Git. In fact, we manage and monitor all of our applications and the whole ‘cloud native stack’ using GitOps. It feels natural and less intimidating to learn, and the tools themselves are very simple.

Git as the Source of Truth

For the last two years, we’ve been running multiple Kubernetes clusters and Prometheus telemetry databases on Amazon Web Services. You can read more about how we provision Kubernetes in the blog post, “[Provisioning And Lifecycle Of A Production Ready Kubernetes Cluster](#)”.

What exactly is GitOps? By using Git as our source of truth, we can operate almost everything. For example, version control, history, peer review, and rollback happen through Git without needing to poke around with tools like kubectl.

- Our provisioning of AWS resources and deployment of k8s is declarative
- Our entire system state is under version control and described in a single Git repository
- Operational changes are made by pull request (plus build & release pipelines)
- Diff tools detect any divergence and notify us via Slack alerts; and sync tools enable convergence
- Rollback and audit logs are also provided via Git



KubeCon



CloudNativeCon

North America 2020

Virtual

Top-level controller implements the management primitives

You write declarative config to implement your operator logic



KubeCon



CloudNativeCon

North America 2020

Virtual

- No longer limited in scale or architecture by needing to fit into the cluster
 - But you will still have to manage RBAC, and handle authentication

When operators are needed



KubeCon



CloudNativeCon

North America 2020

Virtual

...and its state *has* to be managed inside Kubernetes

Definitely not ideal

Engage the app maintainers as early and often as possible

- They might even be willing to take on smaller state-handling feature requests

Your app can take fuller advantage of Kubernetes features

- The operator provides state handling to run your app in older versions of k8s
- The operator provides functions for your app that have not landed in Kubernetes yet

This is fine on a *temporary* basis

You need an operator... but maybe you
shouldn't write it



Virtual

You should not write an operator for things until:

- you've spent as long *recovering from failure or dealing with lack of management*
- as it would reasonably take you to write the operator

Failure is where we learn

- Learning from failure is even more critical when we automate



KubeCon



CloudNativeCon

— North America 2020 —

Virtual

Is that all you're doing?

- If so, look at some of the non-operator methods of handling these tasks first



KubeCon



CloudNativeCon

North America 2020

Virtual

An operator with no forward compatibility assurance is a time sink waiting to happen
Don't assume 1.x means SemVer-level stability is in effect

Case #3: Did you talk to the owner?



KubeCon



CloudNativeCon

— North America 2020 —

Virtual

The app owner knows their app better than anyone
...in fact, they might already be writing what you need

I'm writing an operator, now what?



Virtual

- Understand CRDs and Controllers
- Get started looking at the operator landscape



Virtual

- Can you write in one of your preferred languages?
- Does its model fit both into your environment and into your head?
- Does the framework maintainer provide framework-related services you want to use?

Final thoughts



KubeCon



CloudNativeCon

— North America 2020 —

Virtual

- Maintain loose coupling
- Remember, no code is often better
- Write what you know (then stop)



KubeCon



CloudNativeCon

North America 2020

Virtual

References:

[Devan Goodwin: When Not to Write a Kubernetes Operator](#)

[The New Stack: Kubernetes: When to Use, and When to Avoid, the Operator Pattern](#)

[CoreOS: Introducing Operators: Putting Operational Knowledge into Software](#)

[WeaveWorks: GitOps - Operations by Pull Request](#)

[Kubernetes: Custom Resources, Operator Pattern](#)

Further reading:

[WeaveWorks: What Is GitOps](#)

[Matt Butcher: Is there a Helm and Operators showdown?](#)

THANK YOU!



<https://bitly.com/2HuBJlE+>

